

beyond

PAAS

AWS

FIRECRACKER

SERVERLESS

LAMBDA

MICRONAUT

JAVA

GRAAL VM

QUARKUS

beyond

# Java on AWS Lambda

## AWS Lambda

- Basics (Demo)
- Lambda Execution Environment
- Custom Runtime (Demo)
- AWS Firecracker

## AWS Lambda and Java

- AWS Lambda w/ Java Runtime (Demo)
- Pitfalls
- Accelerating Serverless Java

## Compile Time Frameworks

- Quarkus (Demo)
- Micronaut (Demo)
- Sketchy comparison



# AWS Lambda

## Introduction



# Cloud Computing

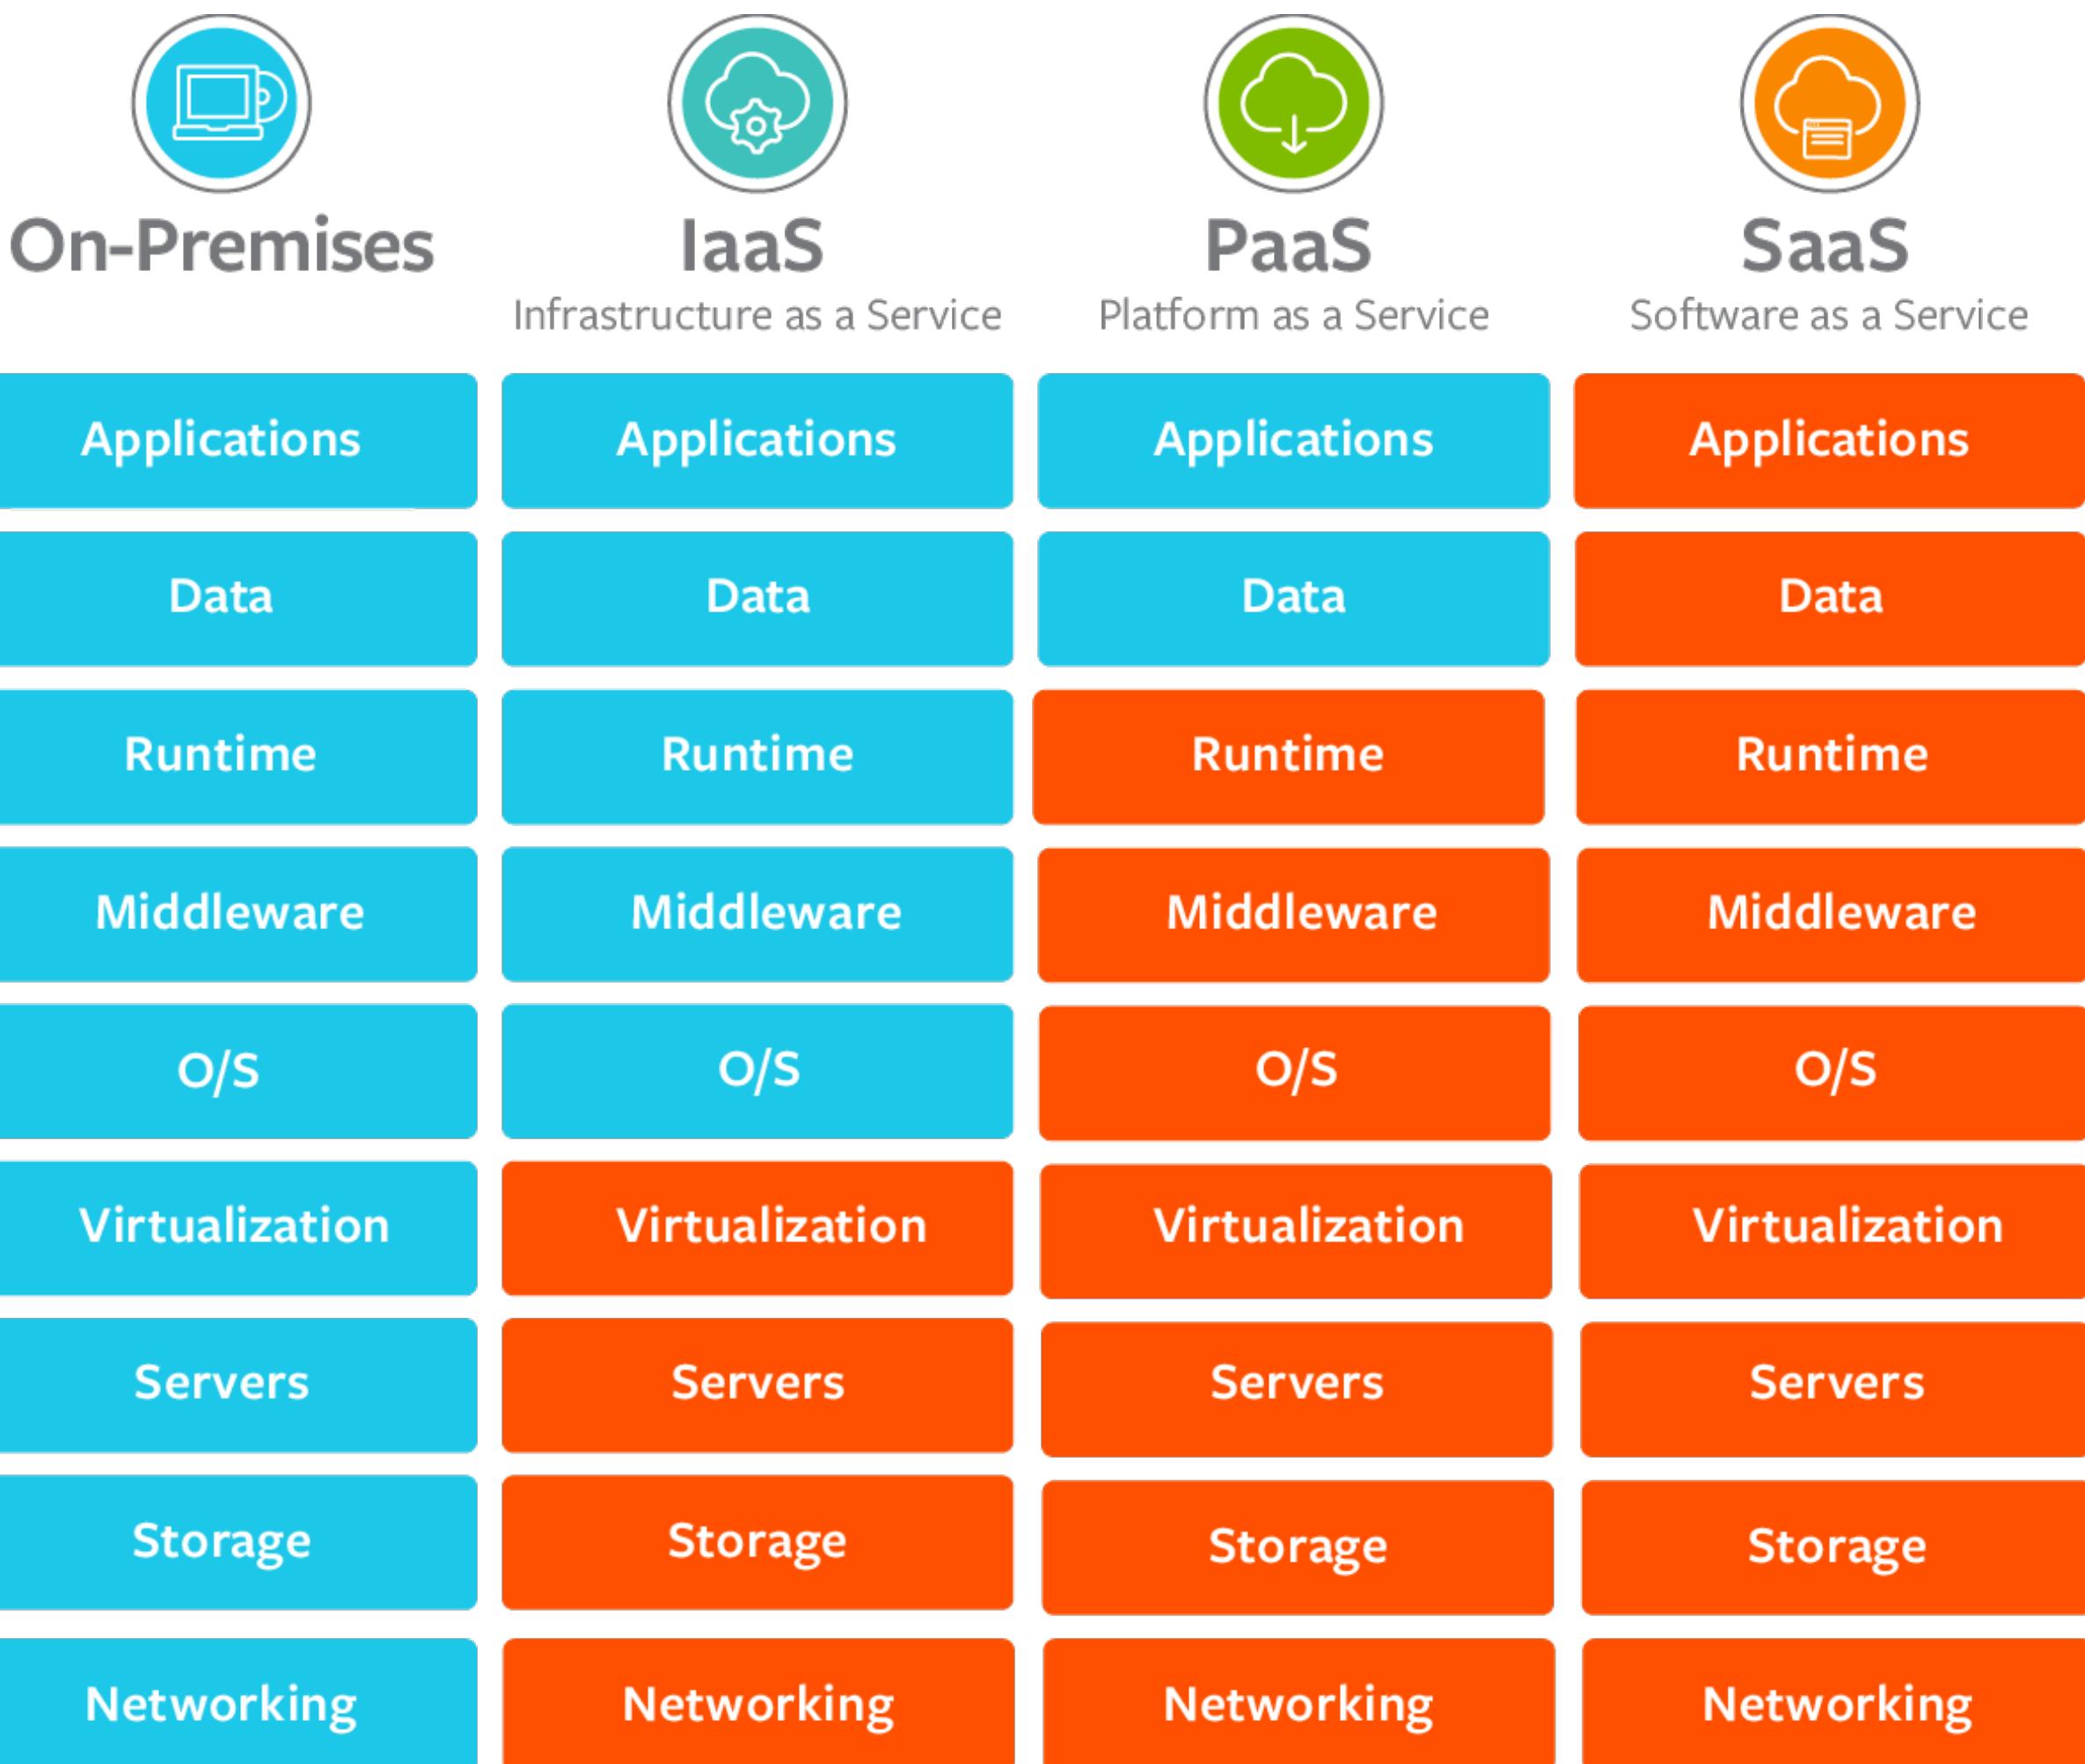
Cloud computing is the **on-demand availability of computer system resources**, especially data storage (cloud storage) and computing power, without direct active management by the user.

## Cloud Computing Service Models

- IaaS
- PaaS
- SaaS

## Cloud Computing Types

- Public cloud is cloud computing that's delivered via the internet and shared across organizations.
- Private cloud is cloud computing that is dedicated solely to your organization.
- Hybrid cloud is any environment that uses both public and private clouds.





# AWS Lambda

## Features

AWS Lambda is a

1. **secure,**
2. **on-demand,**
3. **event-driven,**
4. **compute service**
5. that **scales** as needed (up to 1000 concurrent lambda / account)
6. **billed only for what is used.**

<https://aws.amazon.com/lambda/features/>

## Limitations

- Memory: Max 10 GB
- Max invocation time: **15 minutes**
- Max event size: 256 KB (async) / 6 MB (sync)
- Max function size (zip): 50 MB (zipped) / 250 MB (unzipped)
- Ephemeral storage: 512 MB – 10 GB

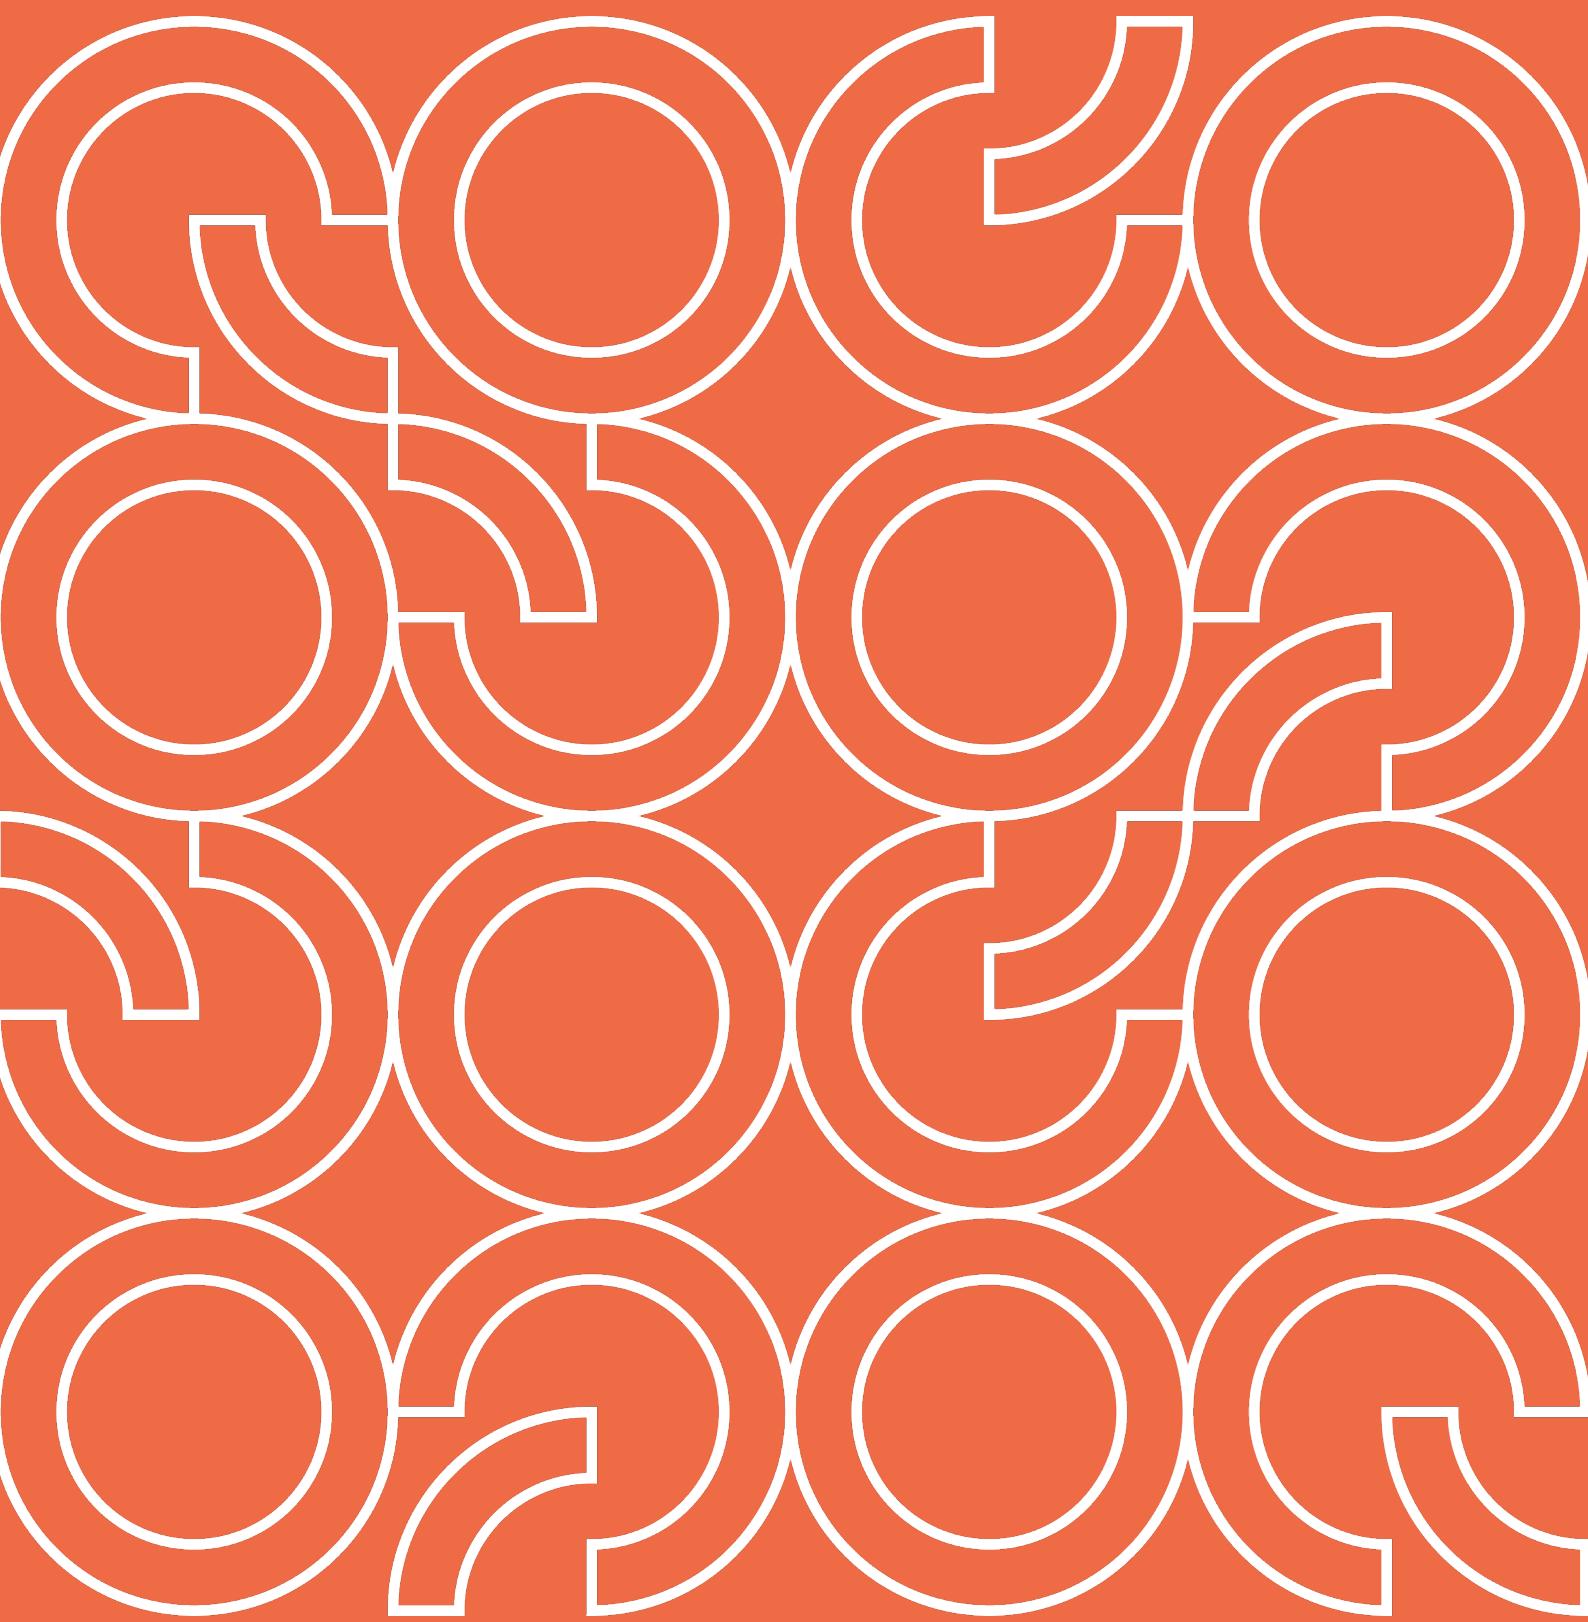
# Demo #1

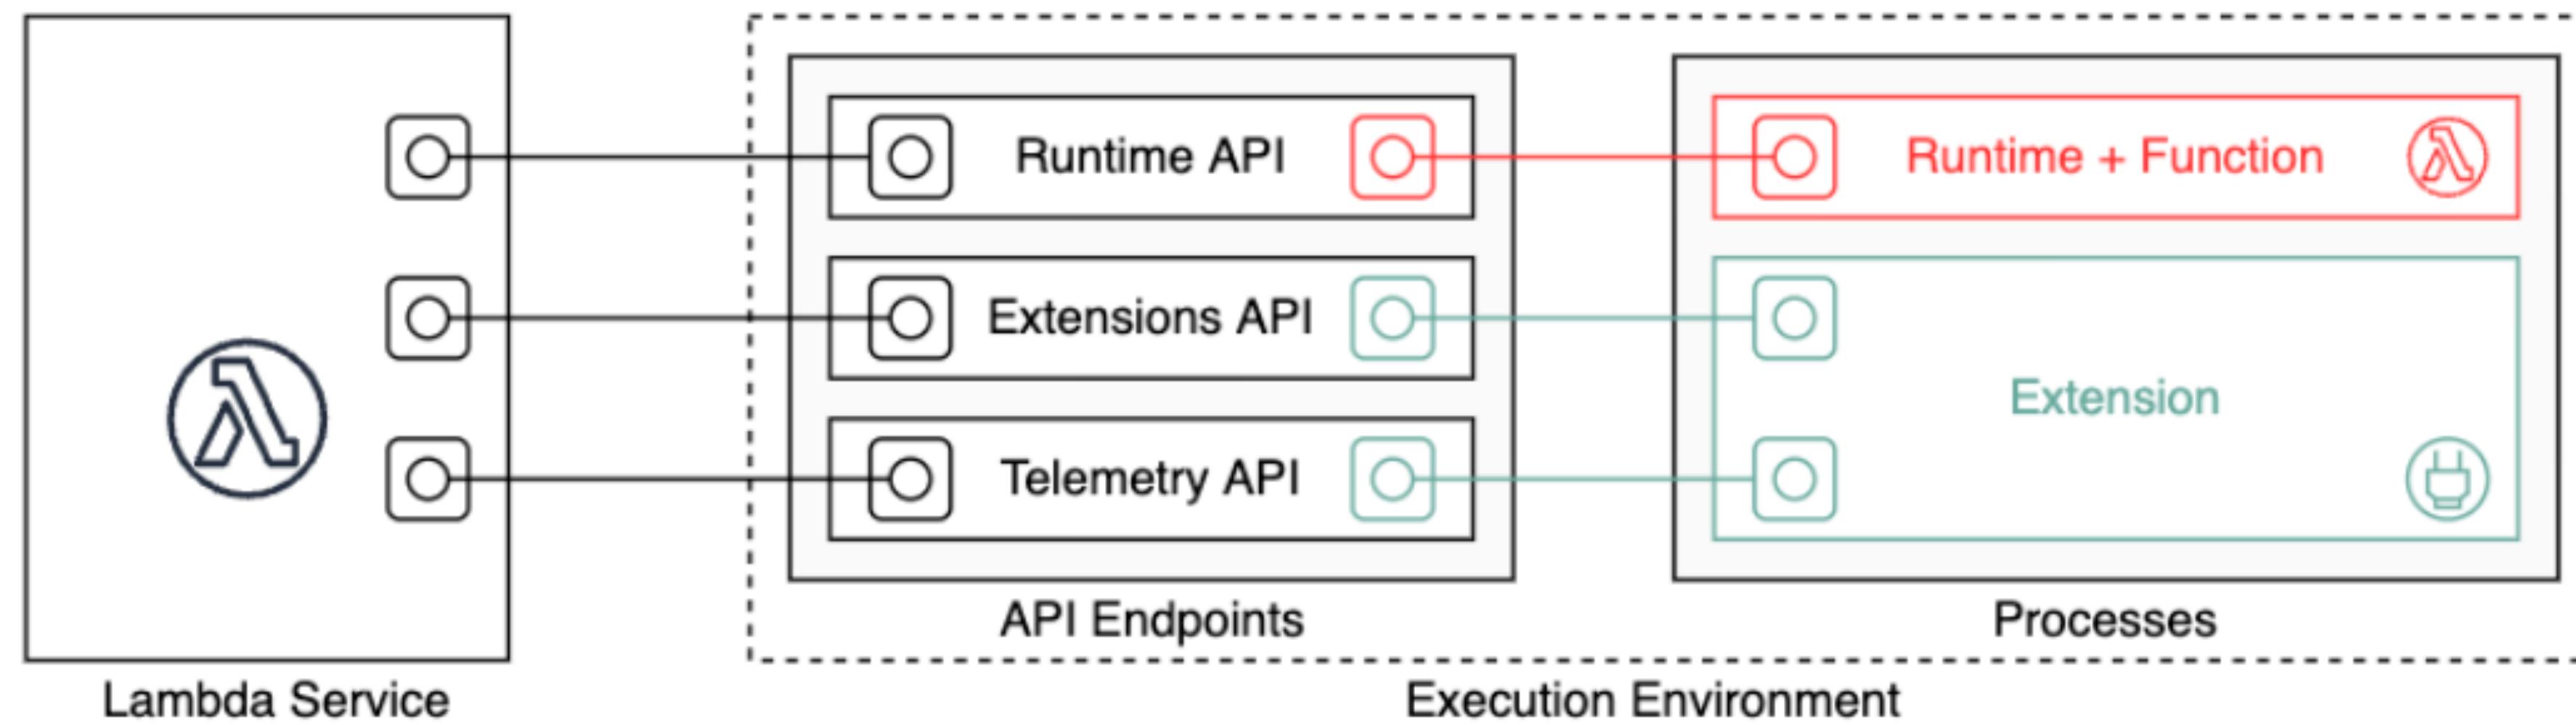
## Lambda basics



# AWS Lambda

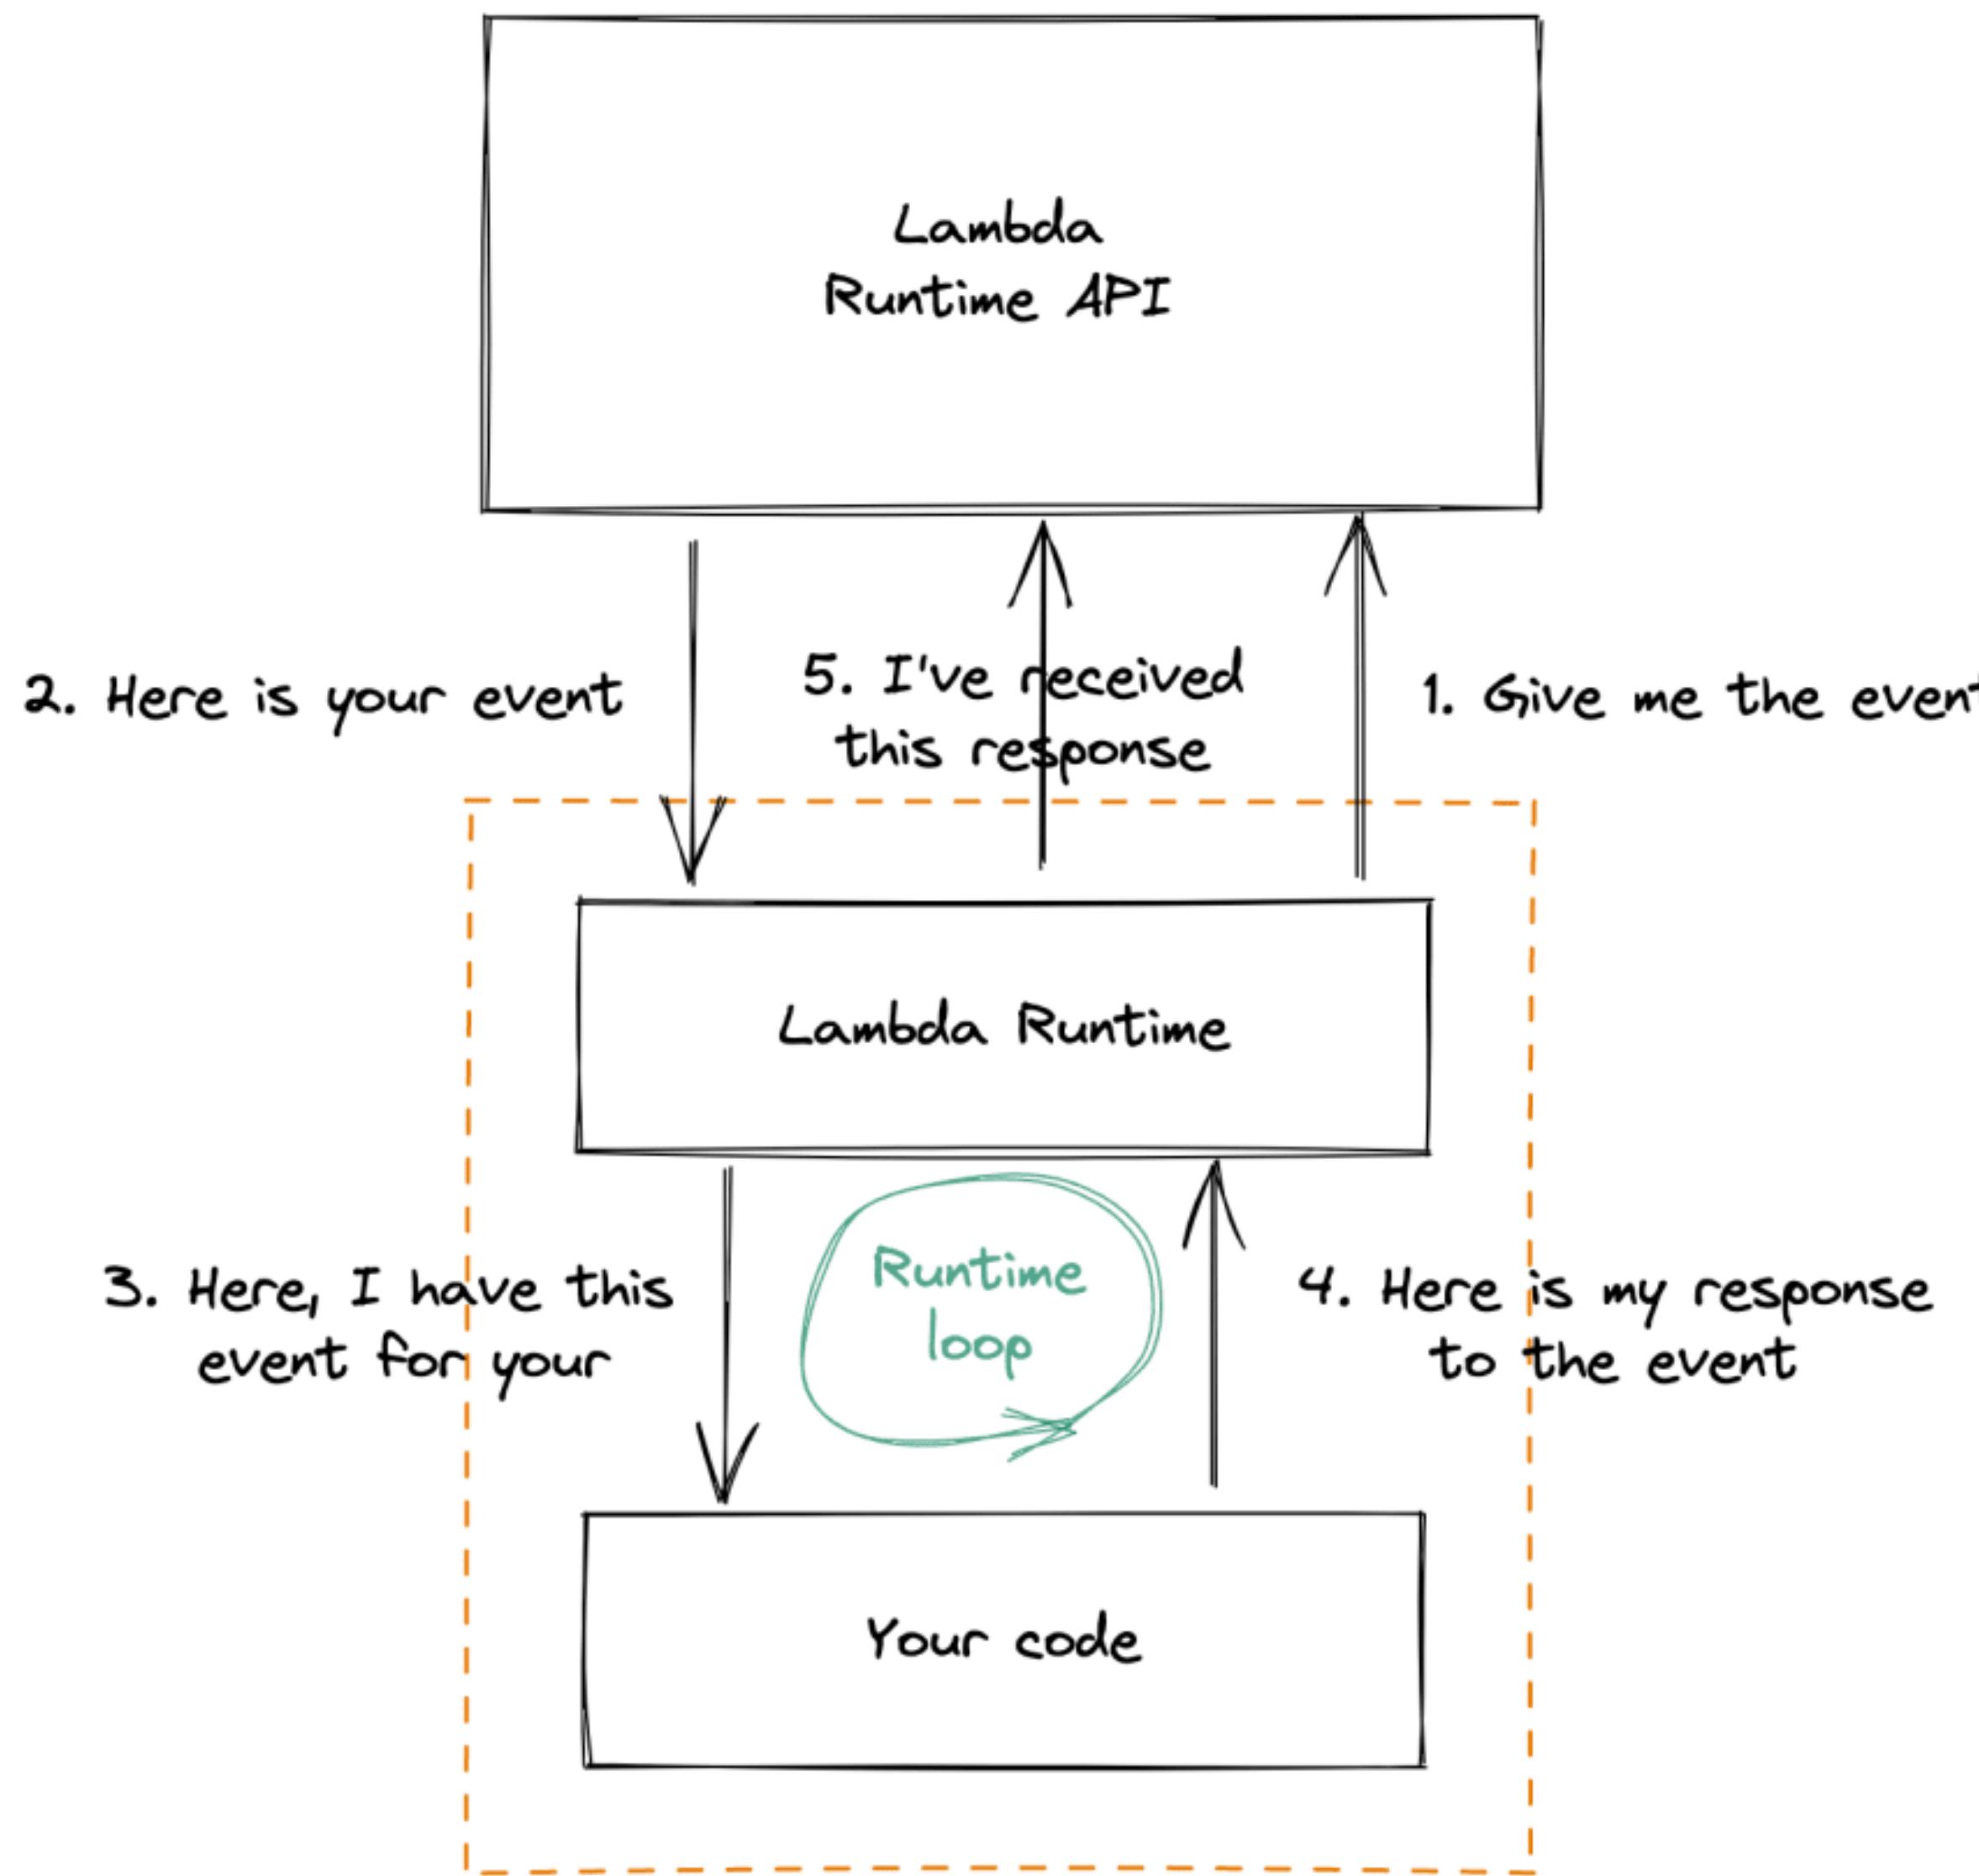
## Under the hood





# Lambda Execution Environment

# Lambda Runtime



A runtime is responsible for running the function's setup code, reading the handler name from an environment variable, and reading invocation events from the Lambda runtime API. The runtime passes the event data to the function handler, and posts the response from the handler back to Lambda.

## Supported Runtimes

- Node.js 18; 16; 14; 12
- Python 3.9; 3.8; 3.7
- Java 8; 11;
- .NET Core 3.1; .NET 5; .NET 6
- Go 1.x
- Ruby 2.7\*
- Custom Runtime

## Operating System

- Amazon Linux v1
- Amazon Linux v2

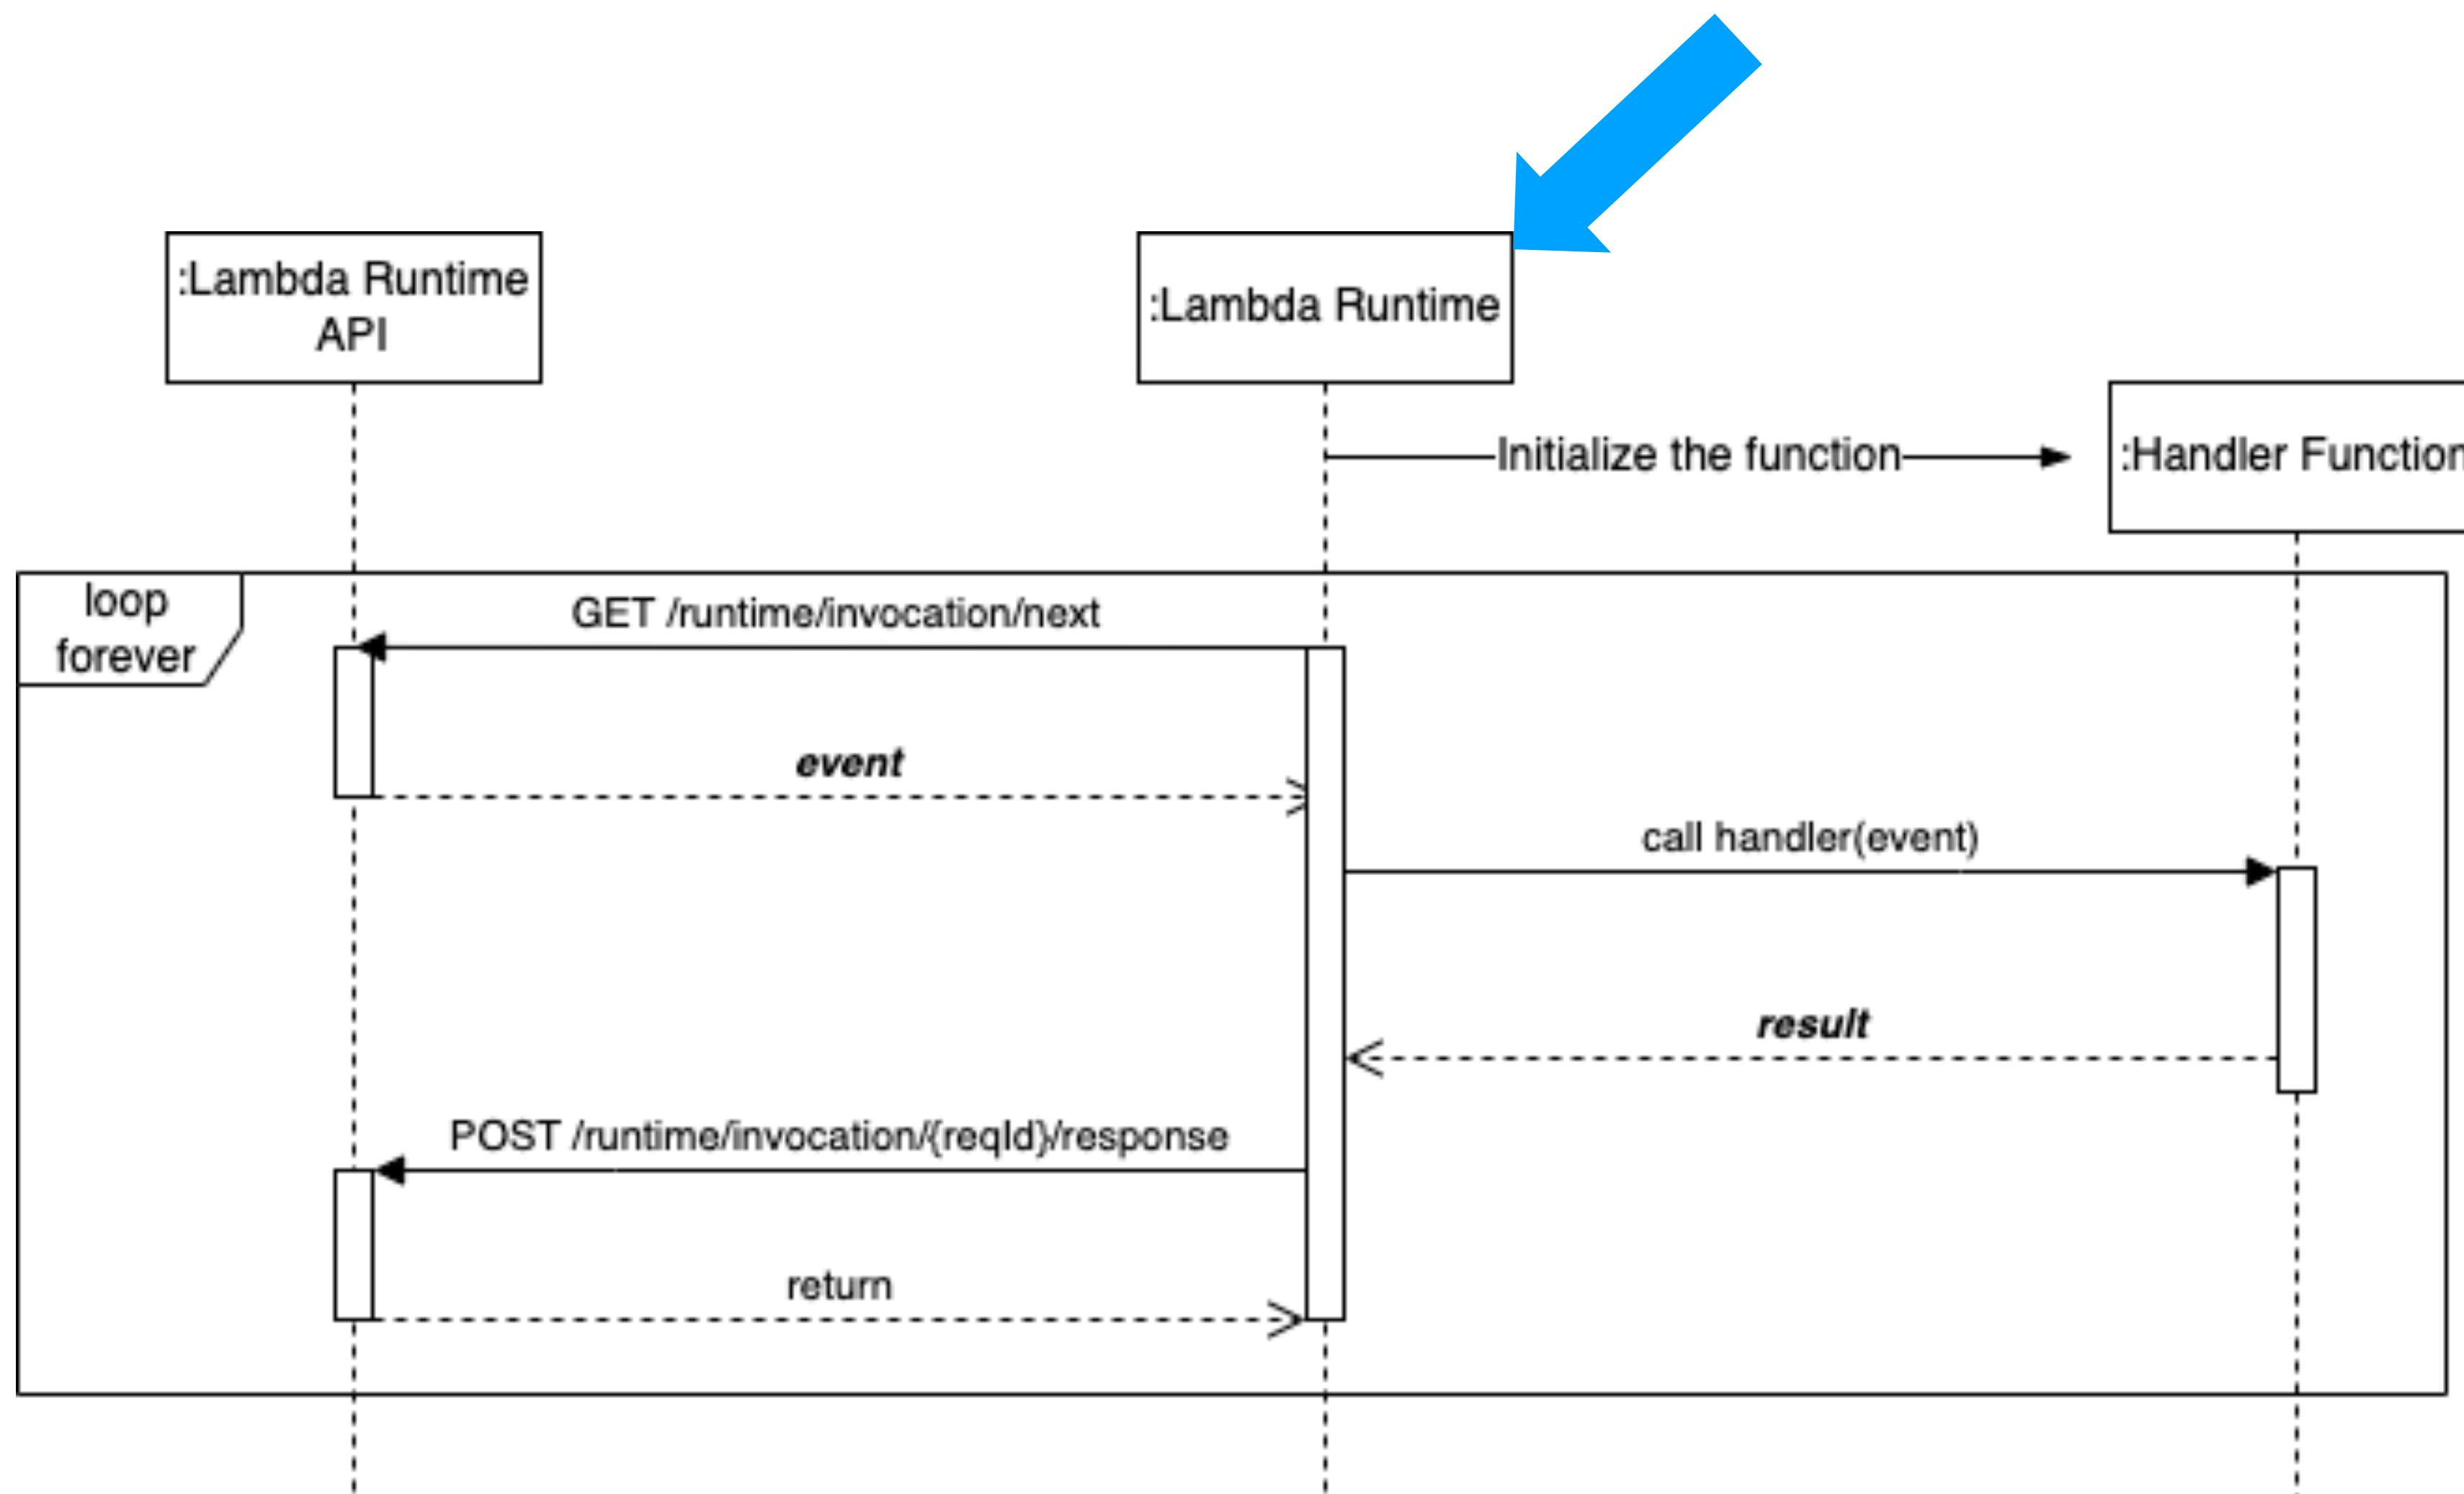


## Runtime API

HTTP API for [runtimes](#) to receive invocation events from Lambda and send response data back within the Lambda [execution environment](#)

### Endpoints

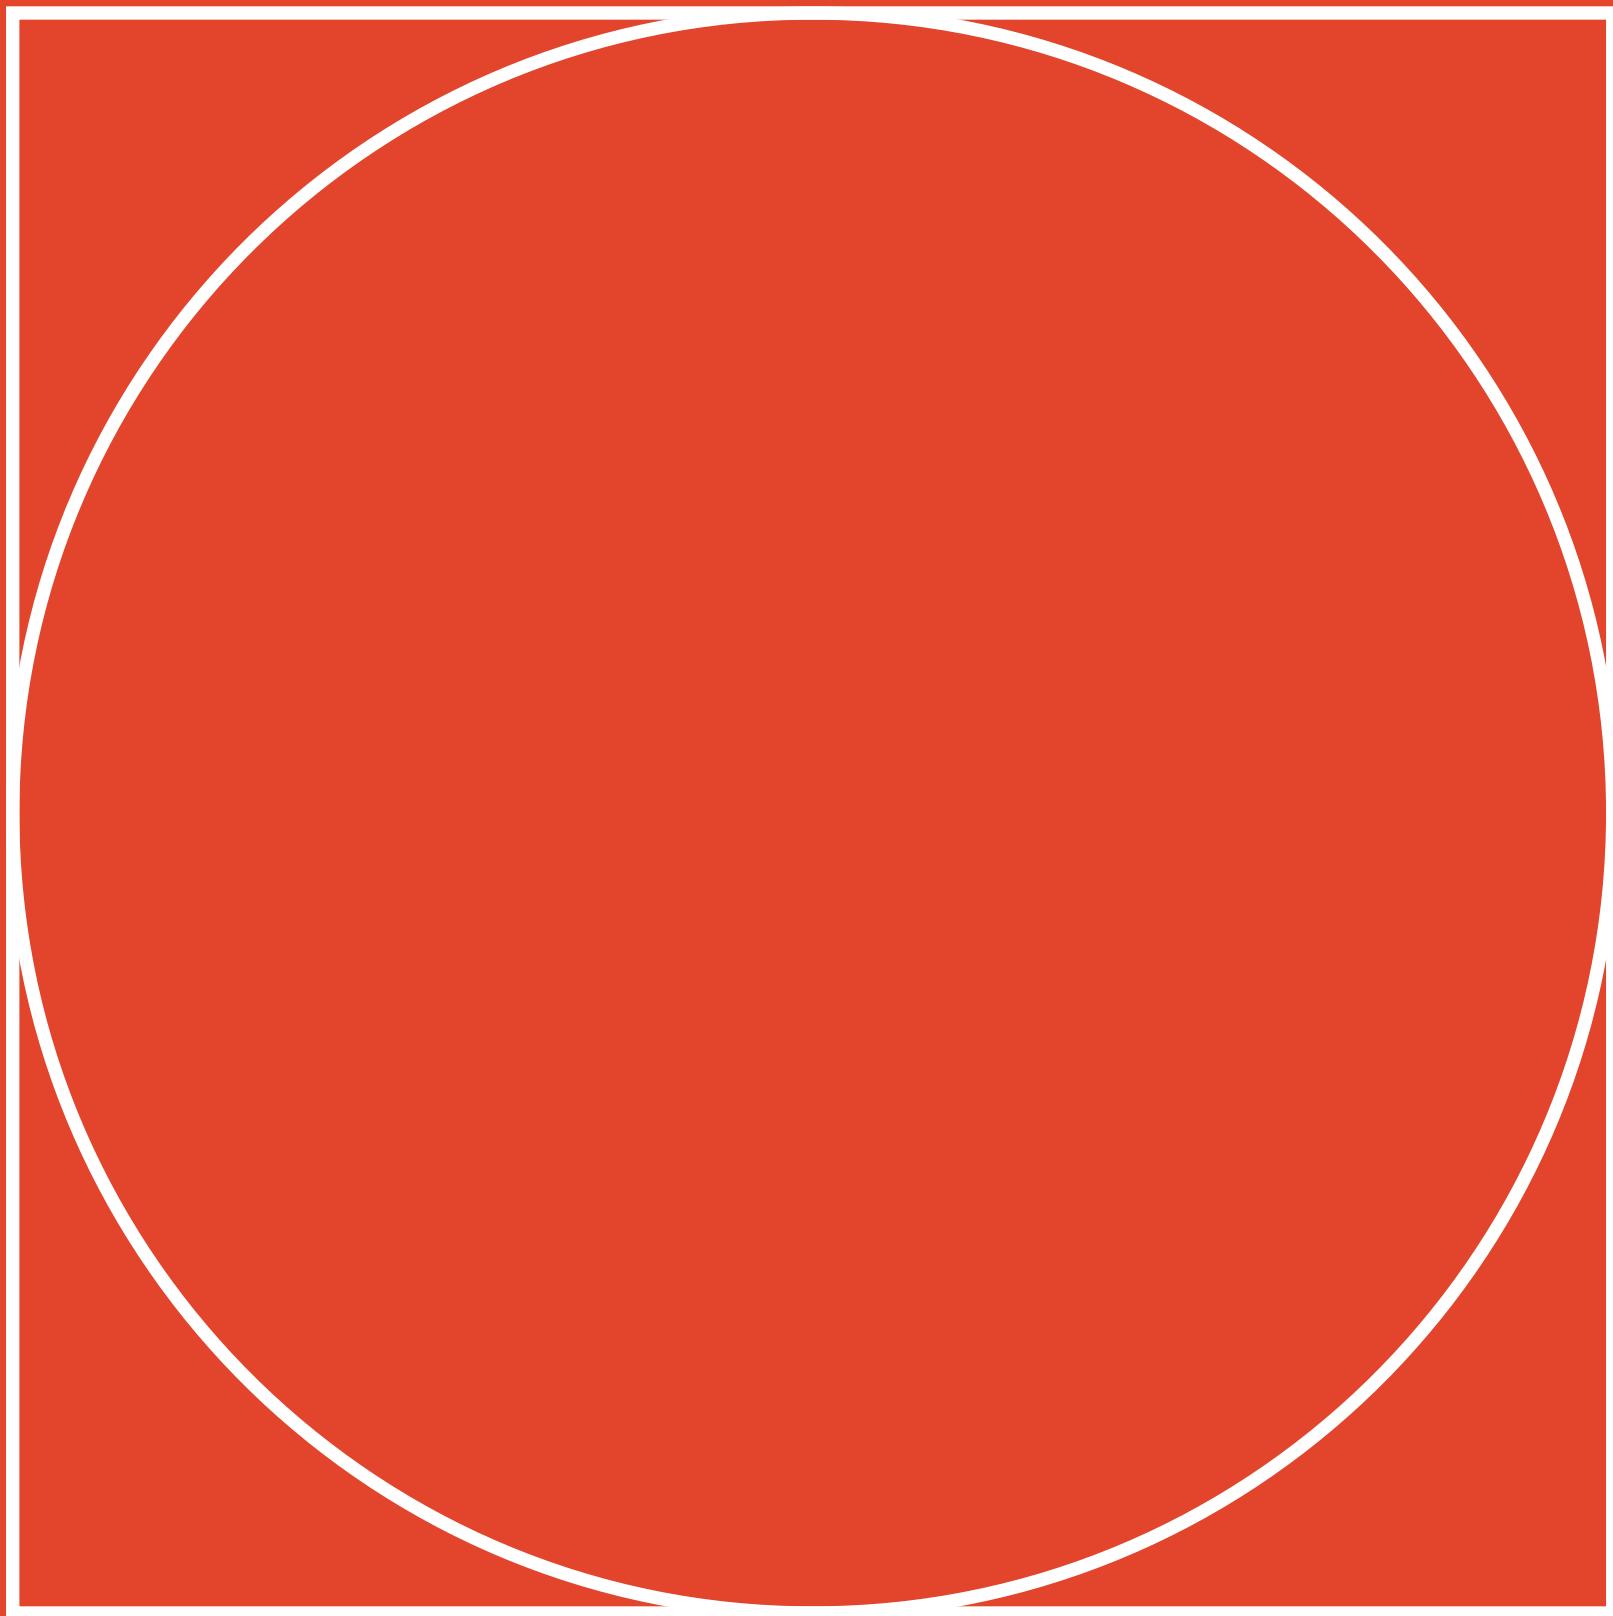
- **GET** /runtime/invocation/next
- **POST** /runtime/invocation/{AwsRequestId}/response
- **POST** /runtime/invocation/{AwsRequestId}/error
- **POST** /runtime/init/error



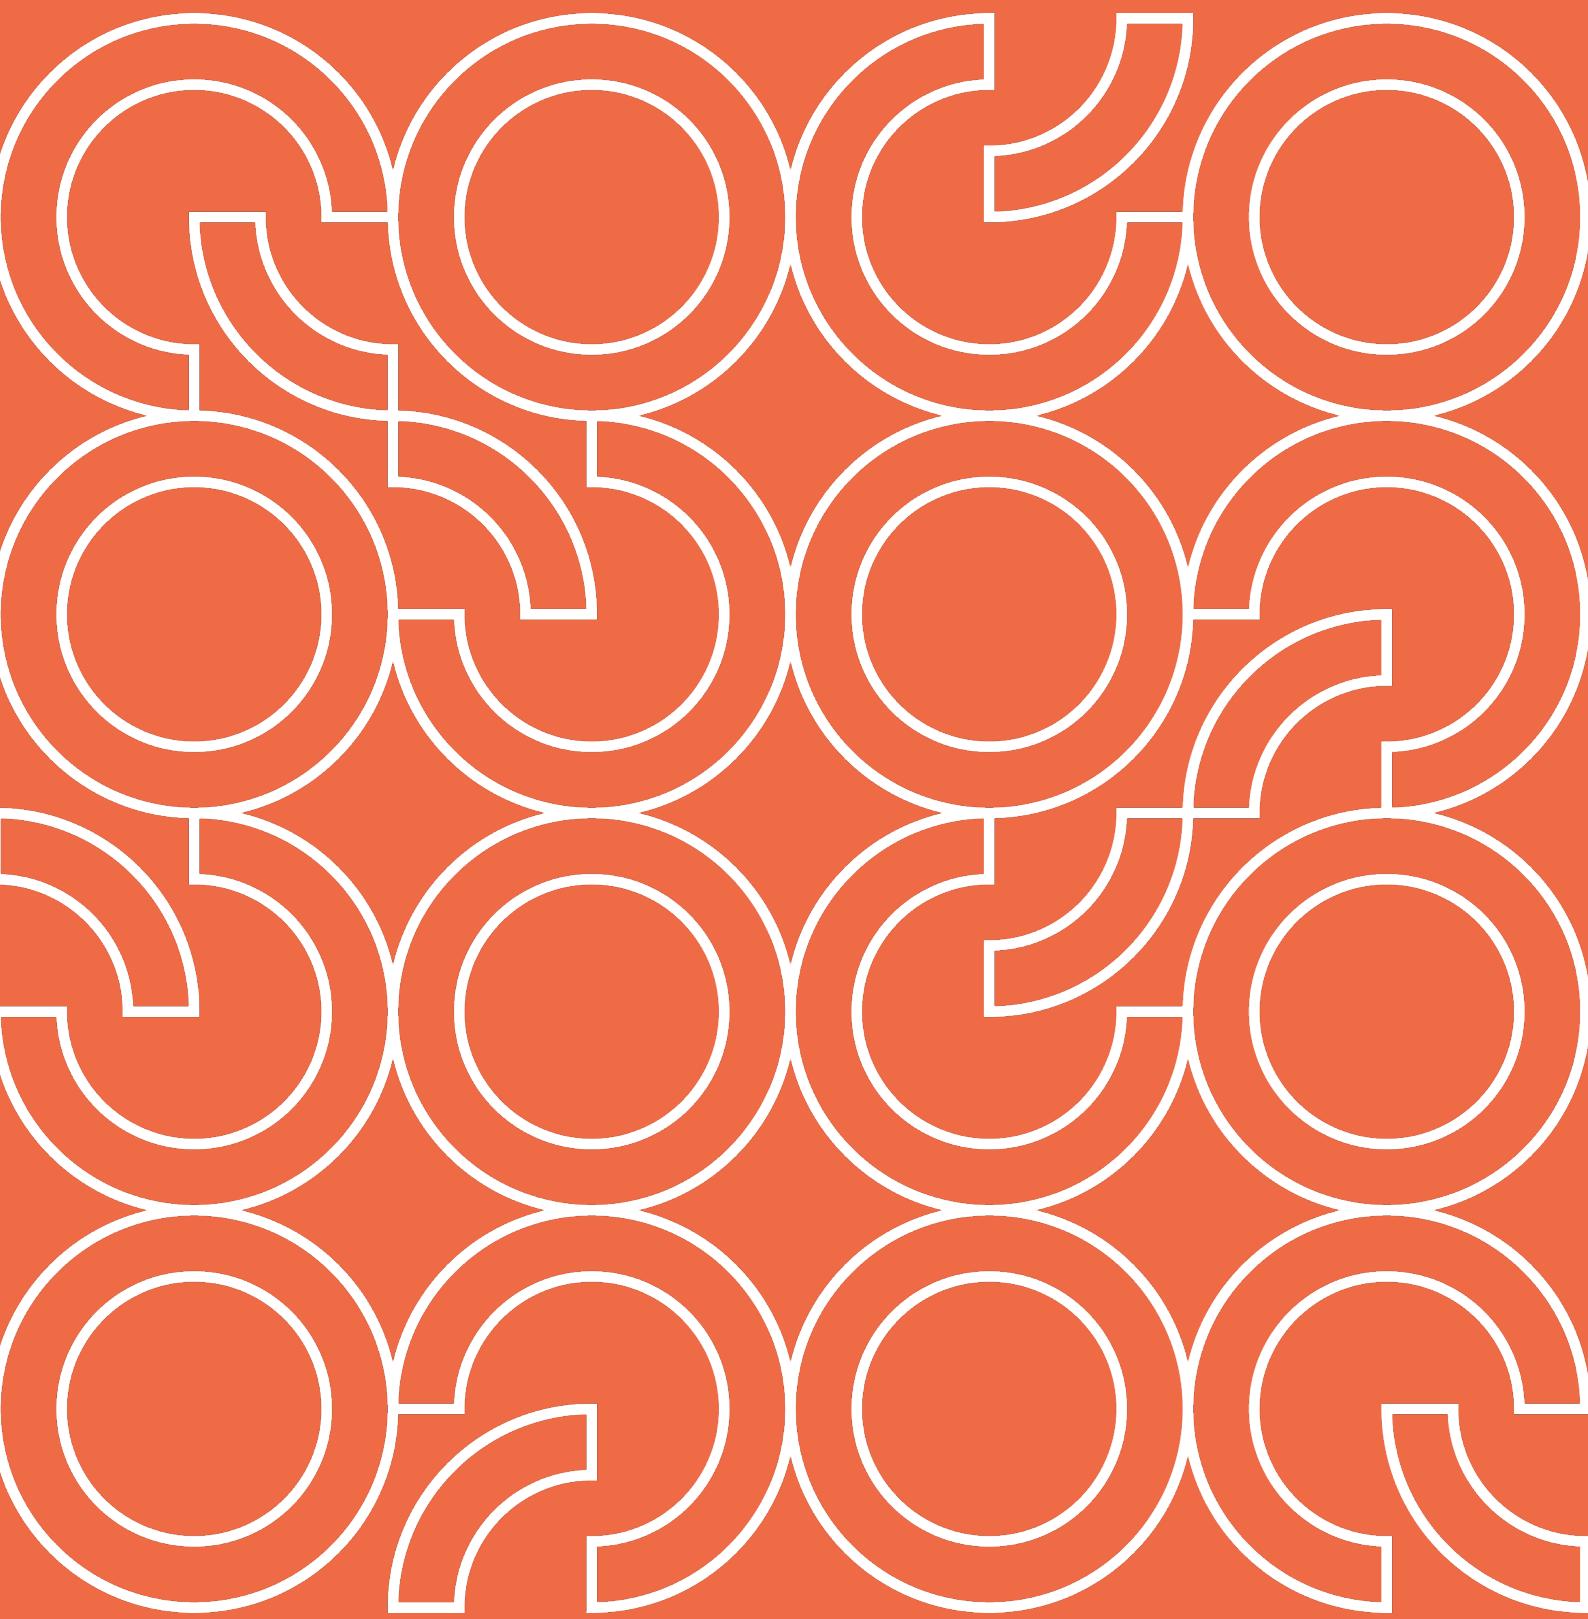
# Implement our own Custom Runtime

## Demo #2

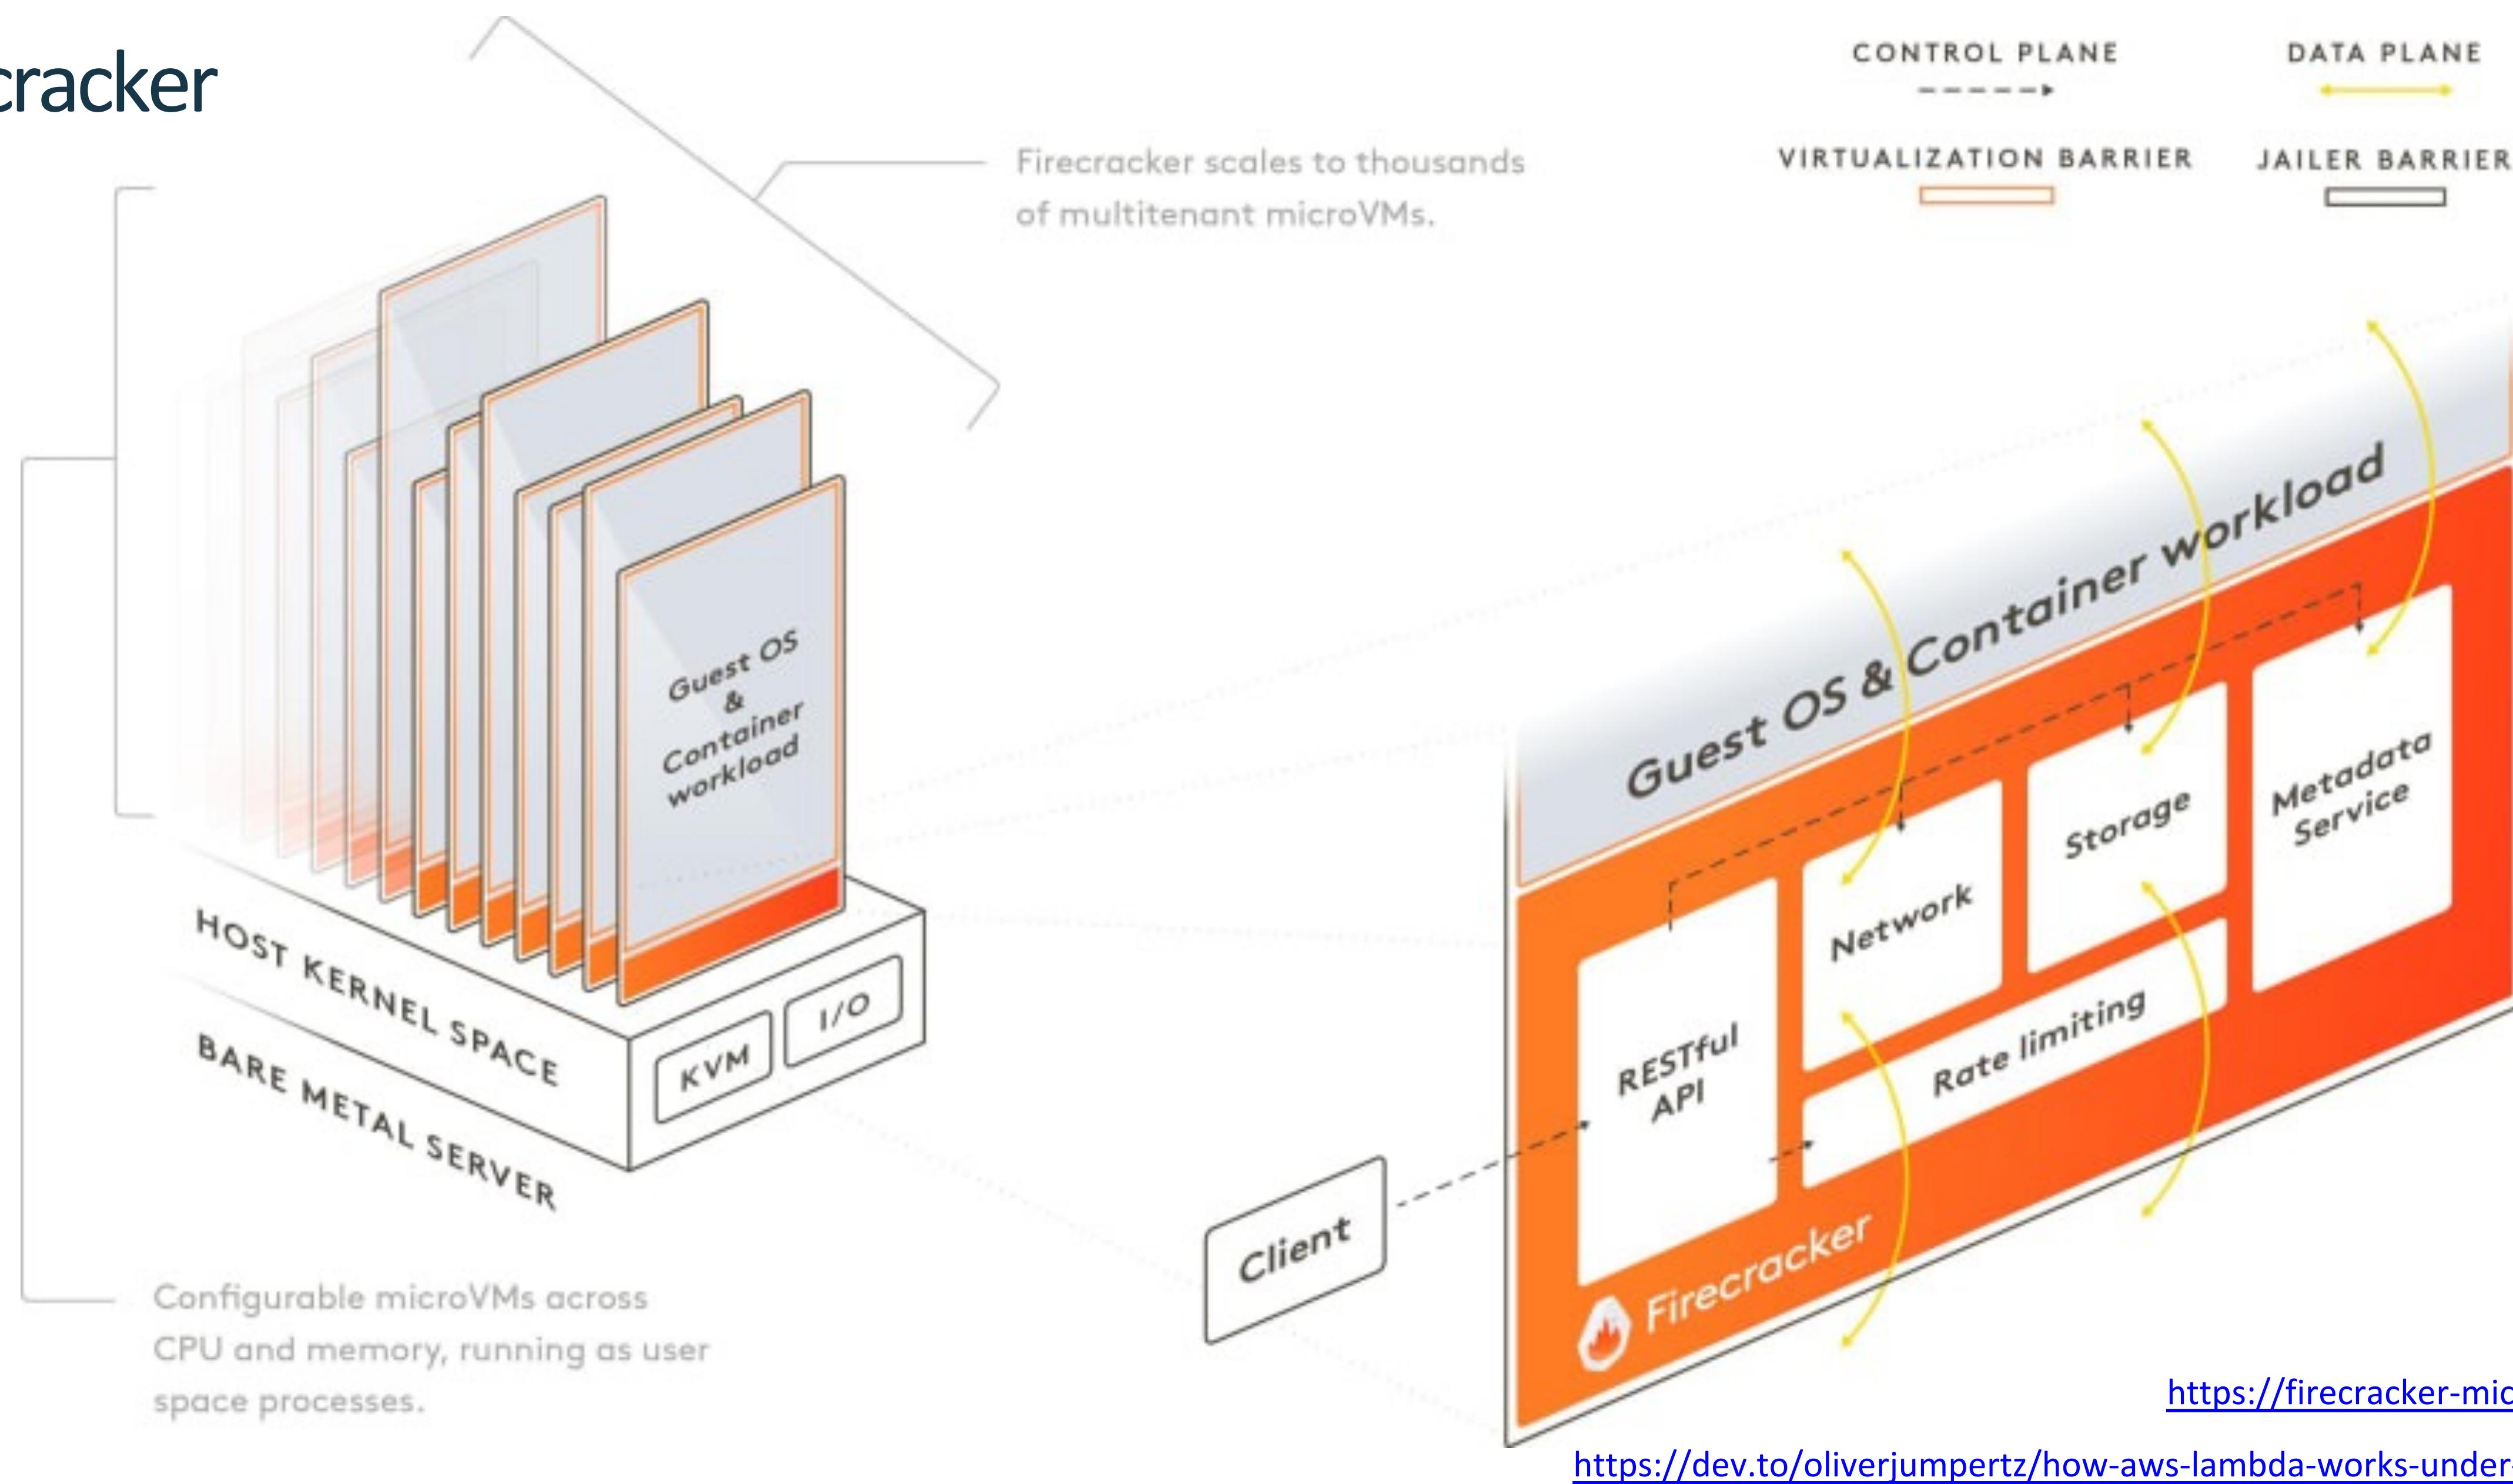
Implement our own Custom  
Runtime



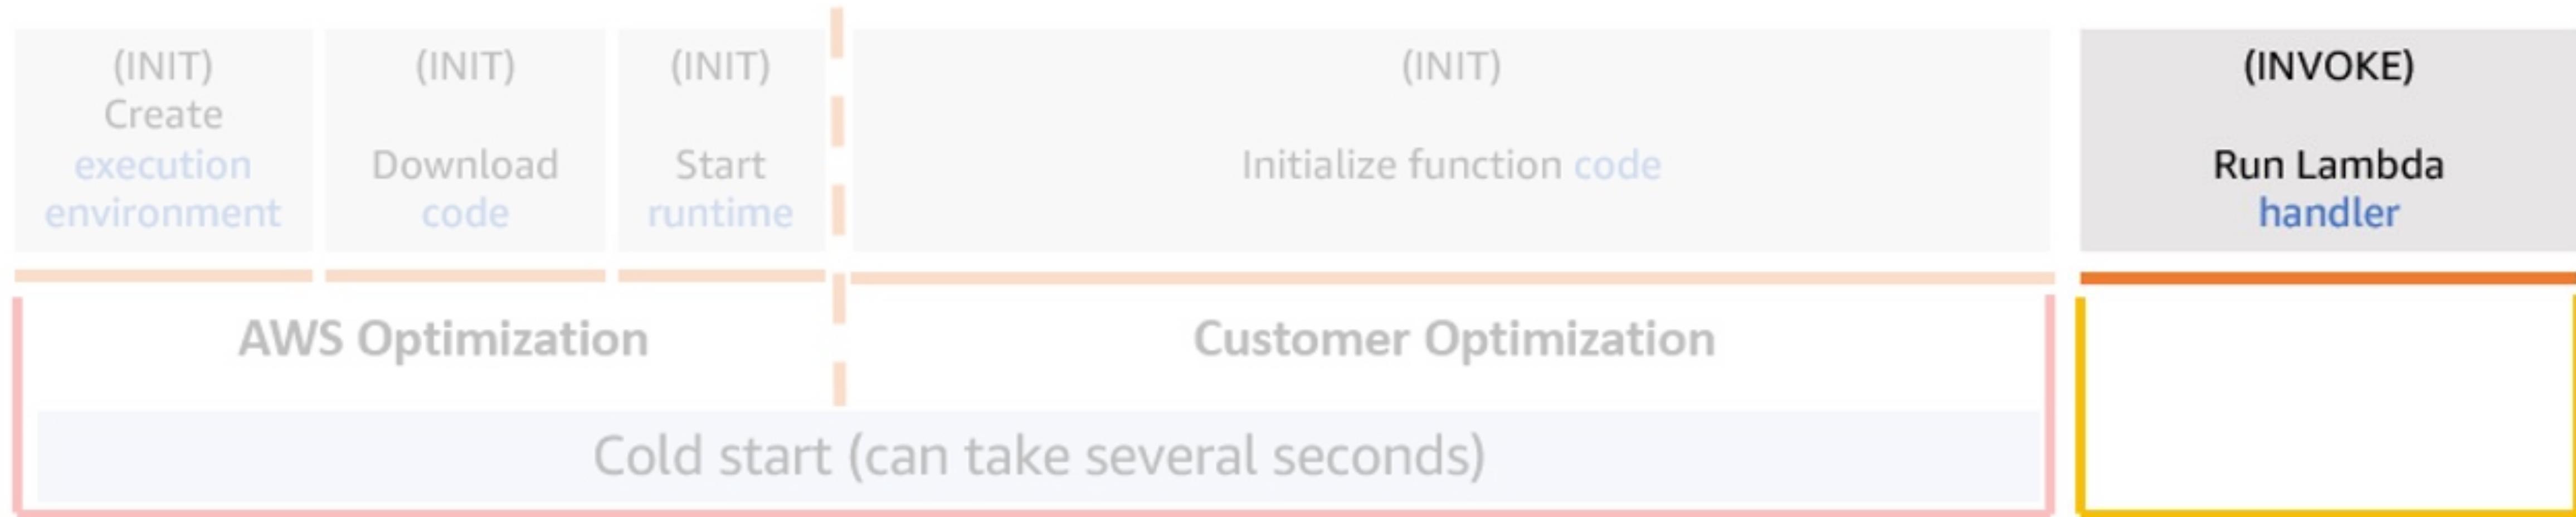
# How AWS Lambda Works?



# Firecracker



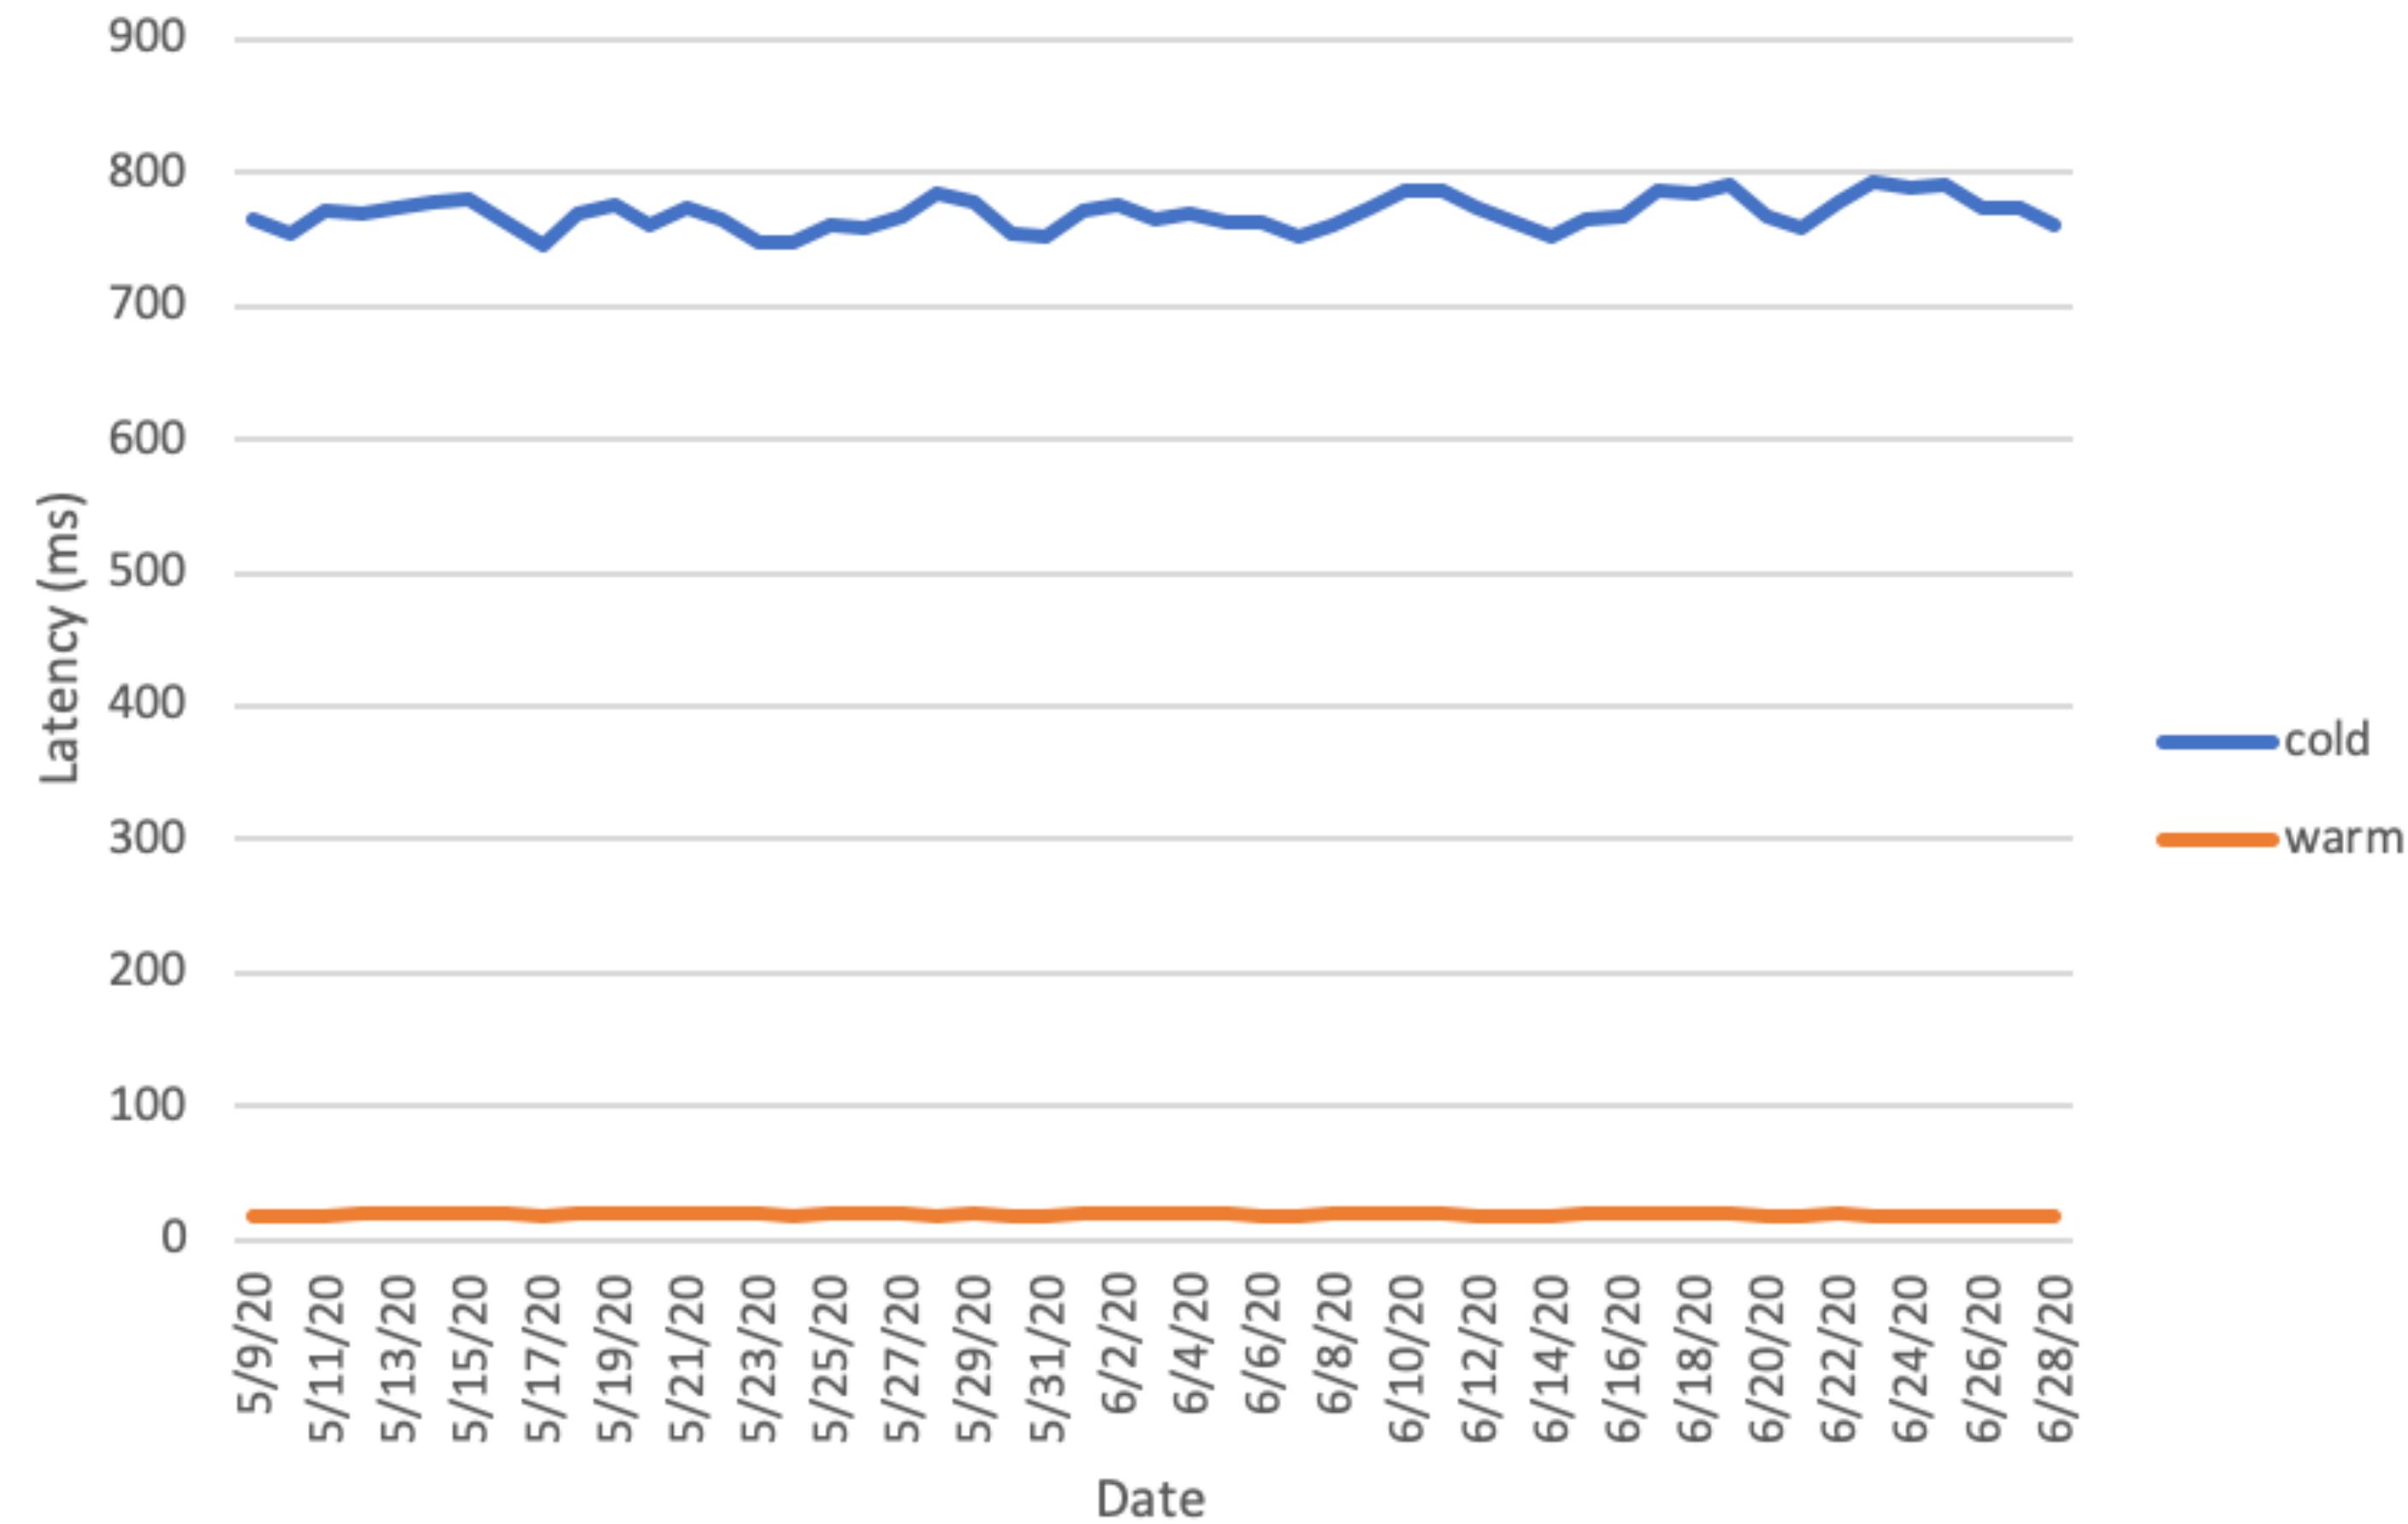
# Cold Start



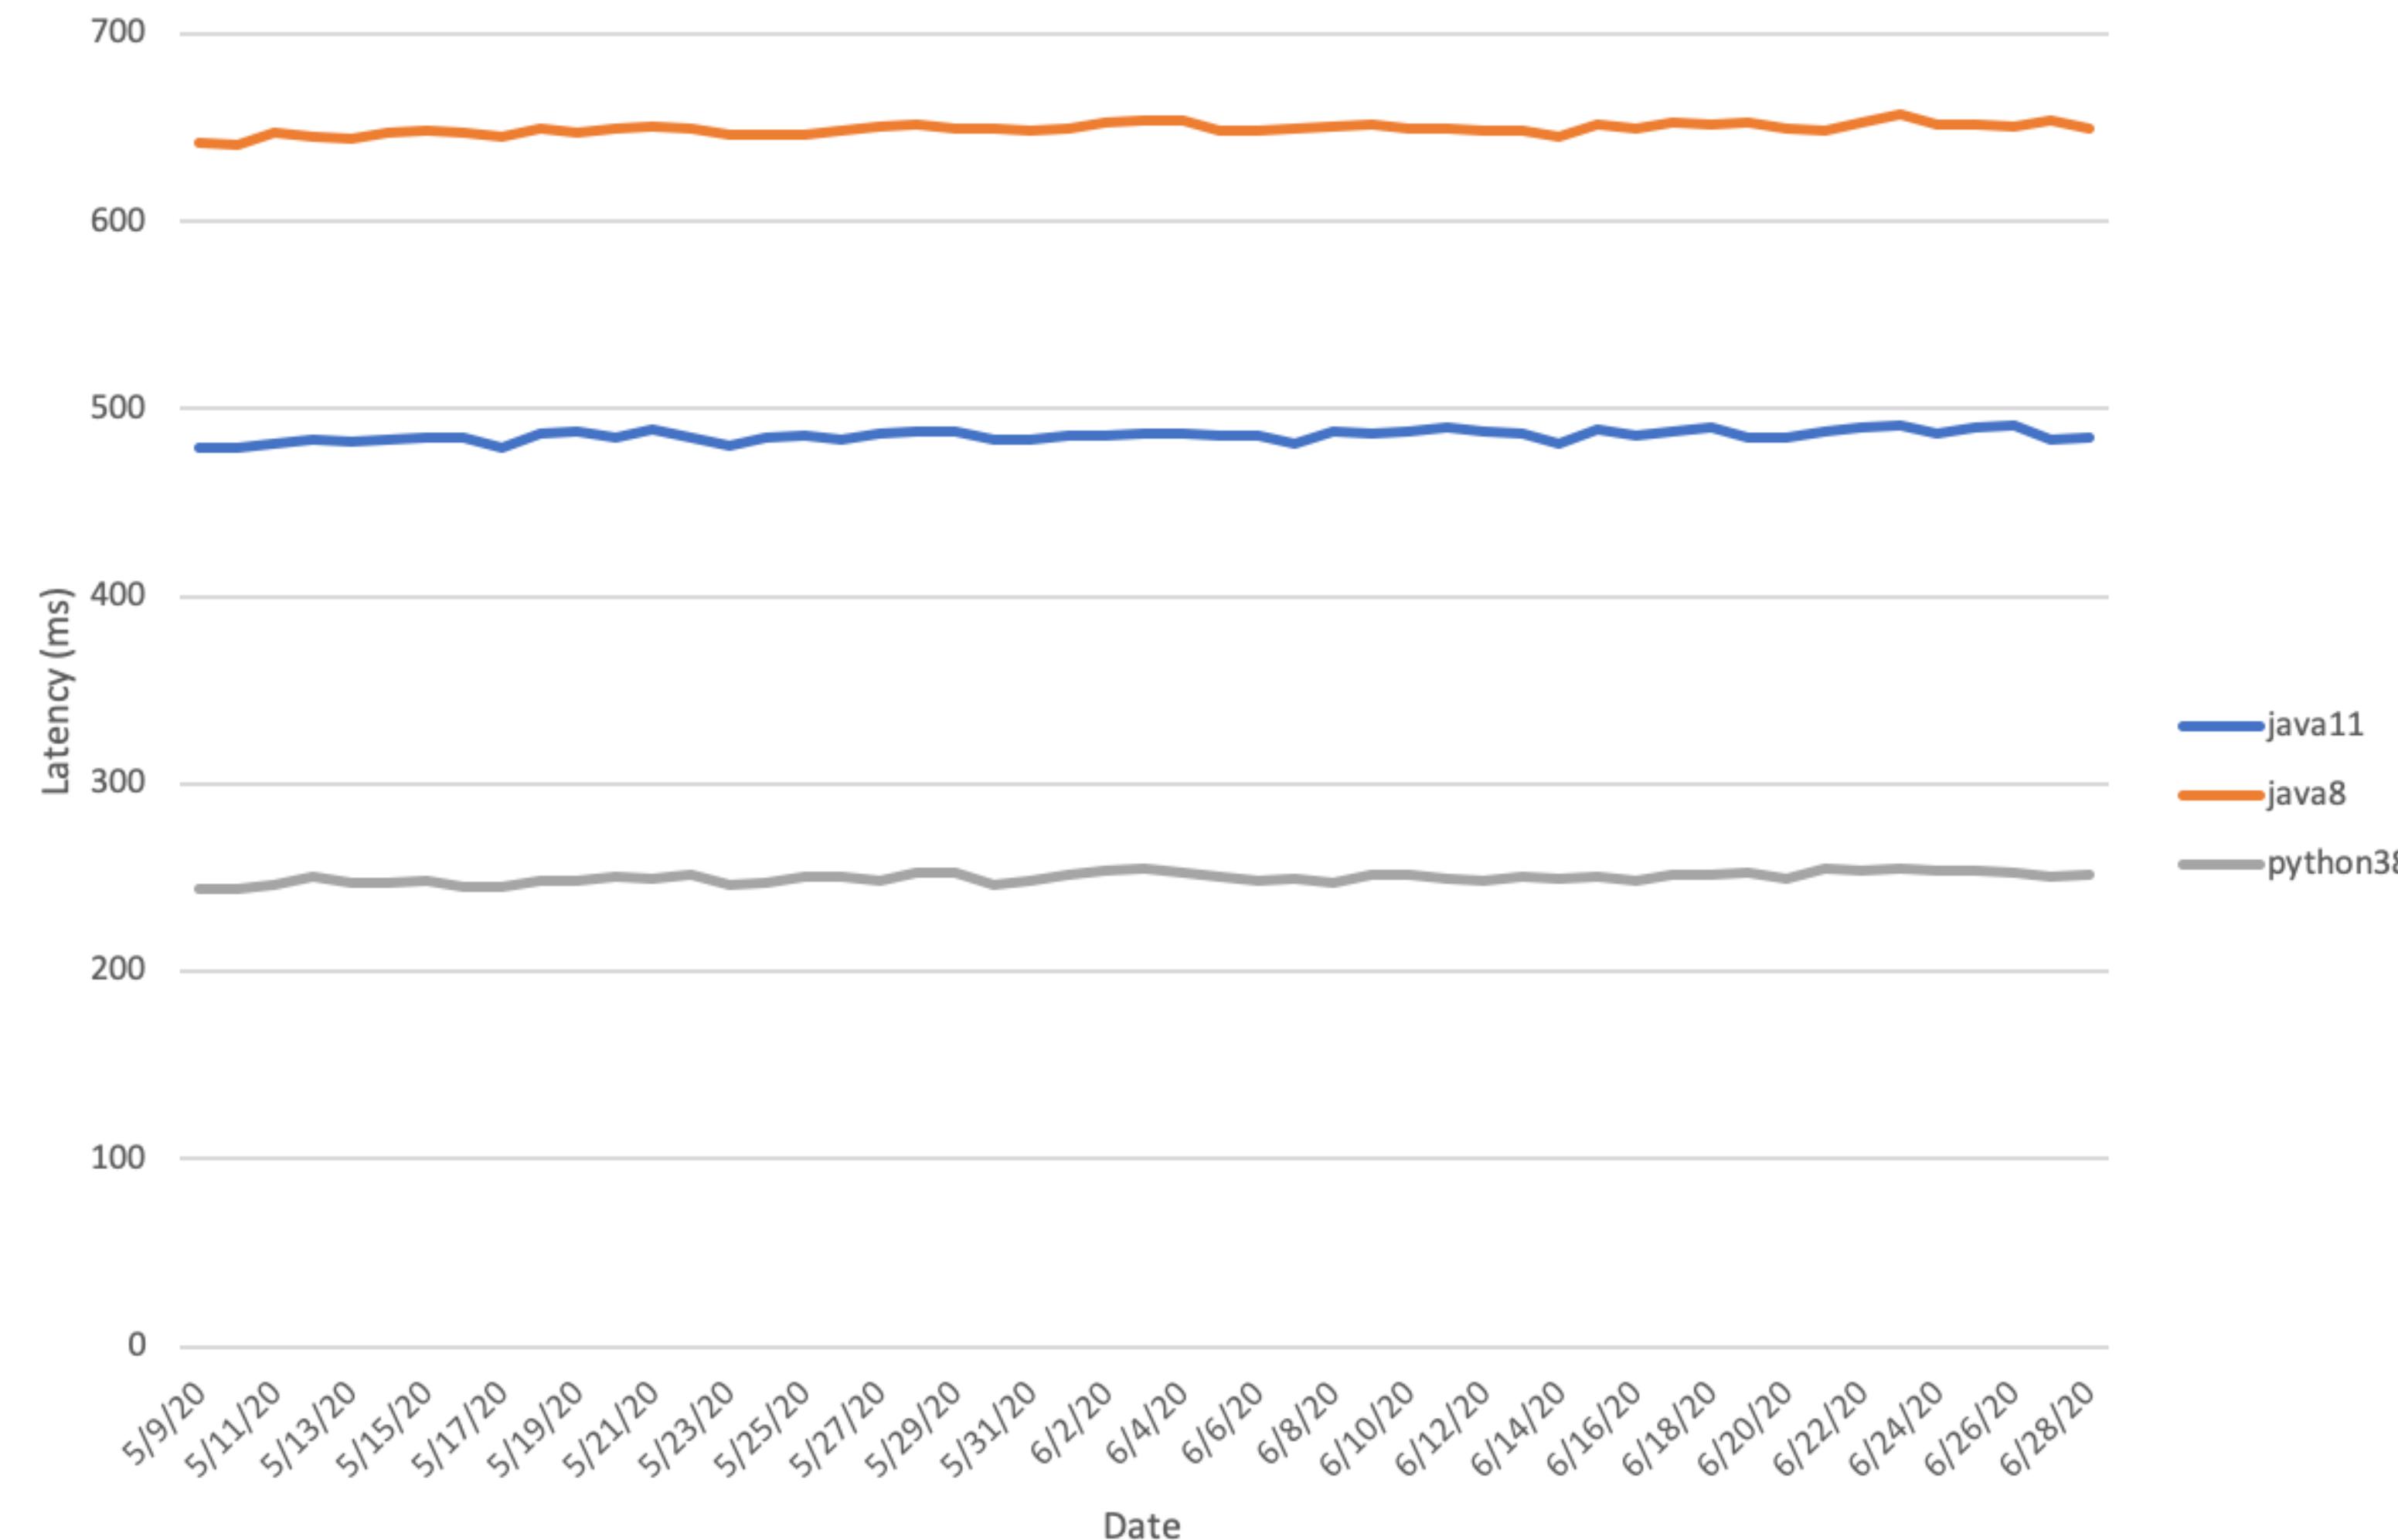
# Warm Start

After the execution completes, the execution environment is frozen. To improve resource management and performance, the Lambda service retains the execution environment for a non-deterministic period of time. During this time, if another request arrives for the same function, the service may reuse the environment. This second request typically finishes more quickly, since the execution environment already exists and it's not necessary to download the code and run the initialization code. This is called a “warm start”.

# Java - Cold Start vs. Warm Start

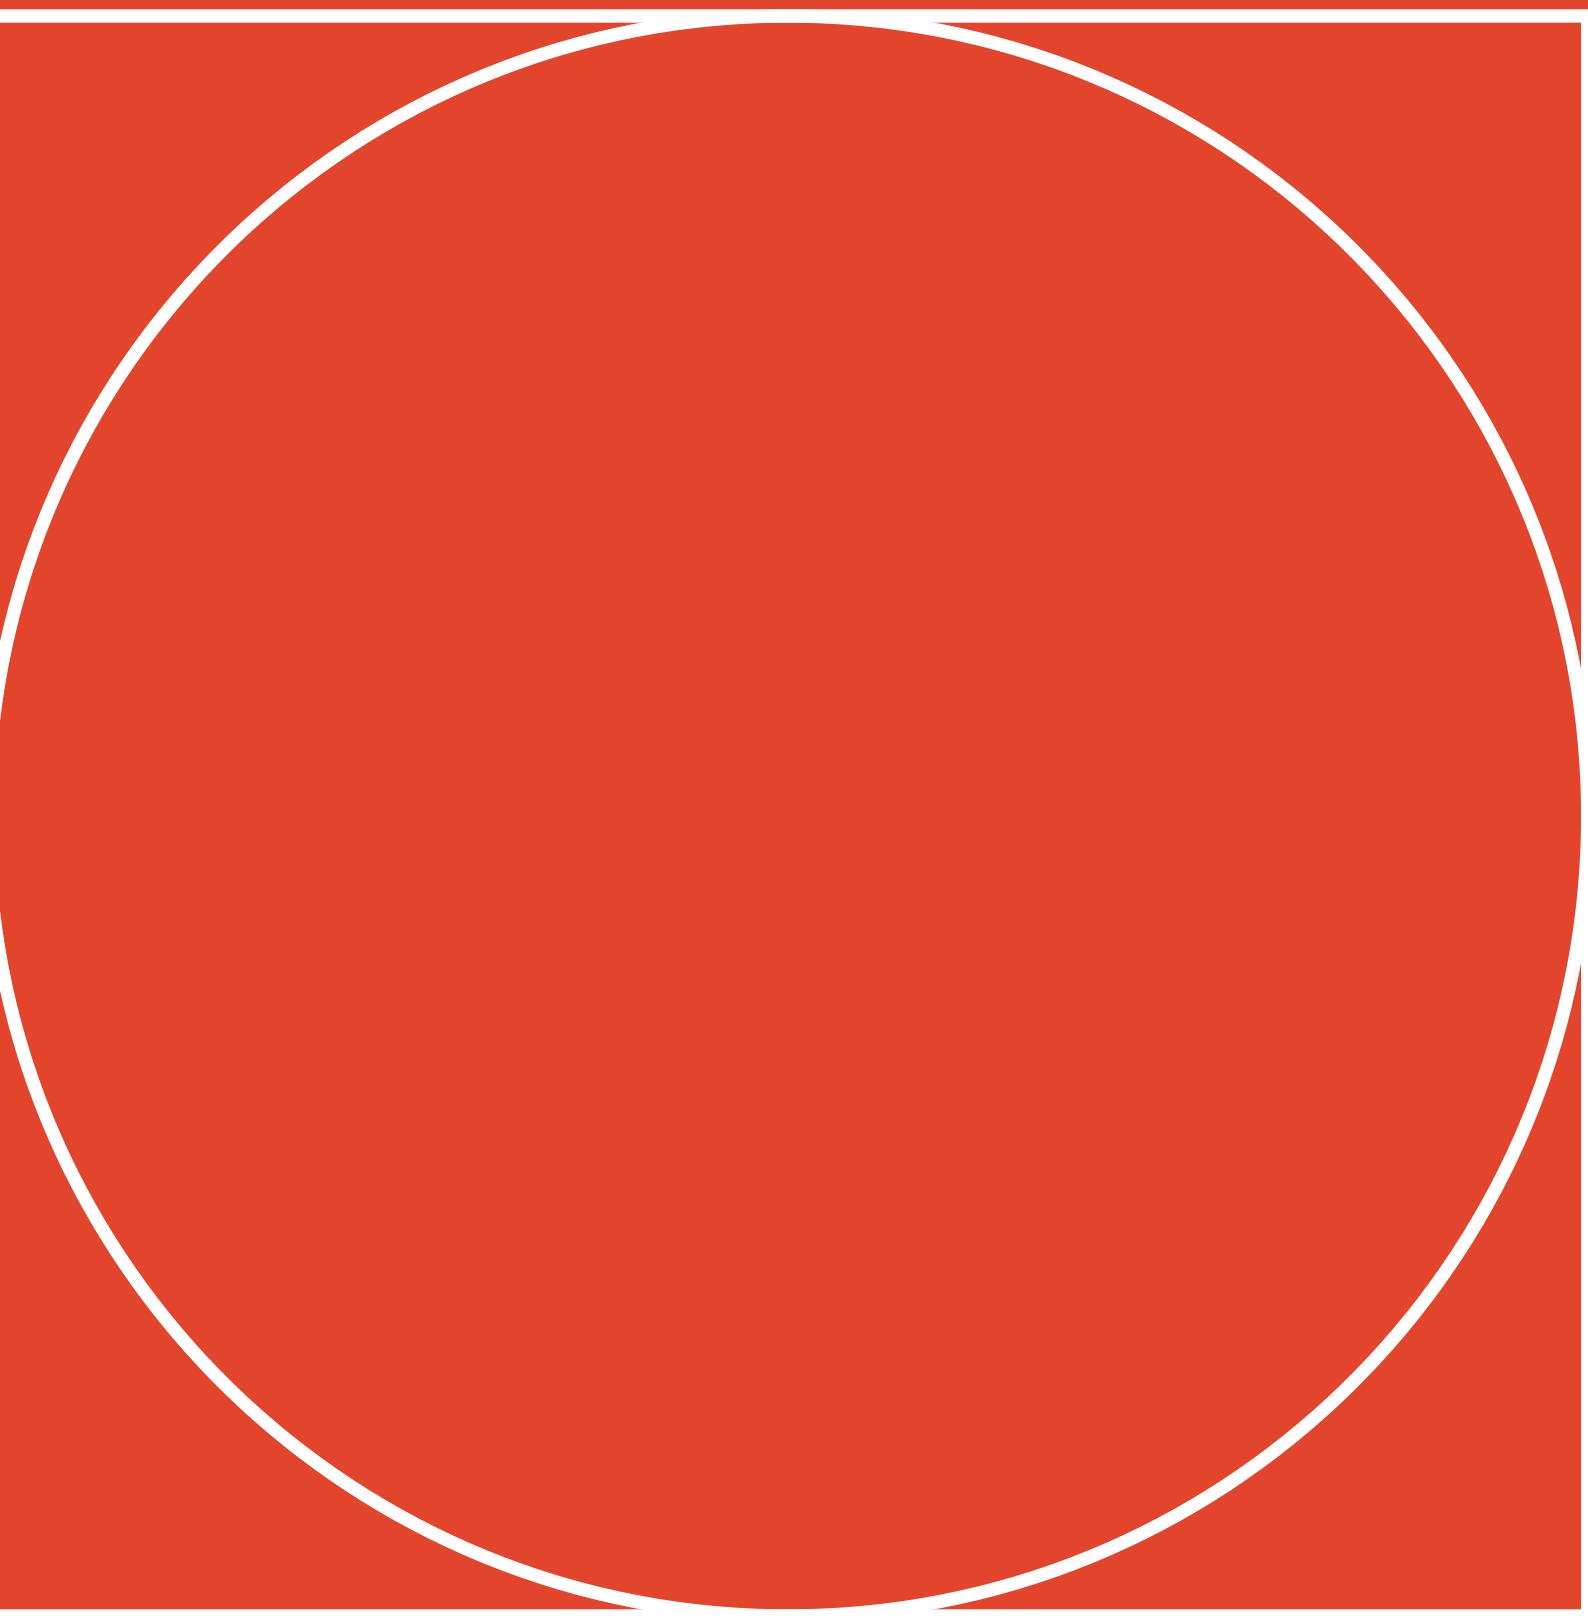


# Java vs. Other Runtime Cold Start



# Demo #3.1

## AWS Lambda w/ Java Runtime





# Accelerate Serverless Java

## Keep it warm

Ping Lambda function in every 5-15 minutes thus AWS Lambda Service won't destroy the execution environment (i.e.: firecracker microVM)

### Disadvantages

- Extra mechanism required (i.e.: additional functionality in lambda, scheduled trigger)
- Does not solve the issue in case of unpredictable number of lambda instances (scaling out)
- Dirty workaround

"I ping my lambda function every 15 minutes to make sure it's always warm"



# Accelerate Serverless Java

## Tiered Compilation

- Use JVM C1 compiler only (disable C2 and the required profiling)
- The C1 compiler is optimized for fast start-up time
- ~40% improvement

JAVA\_TOOL\_OPTIONS to -XX:+TieredCompilation -  
XX:TieredStopAtLevel=1

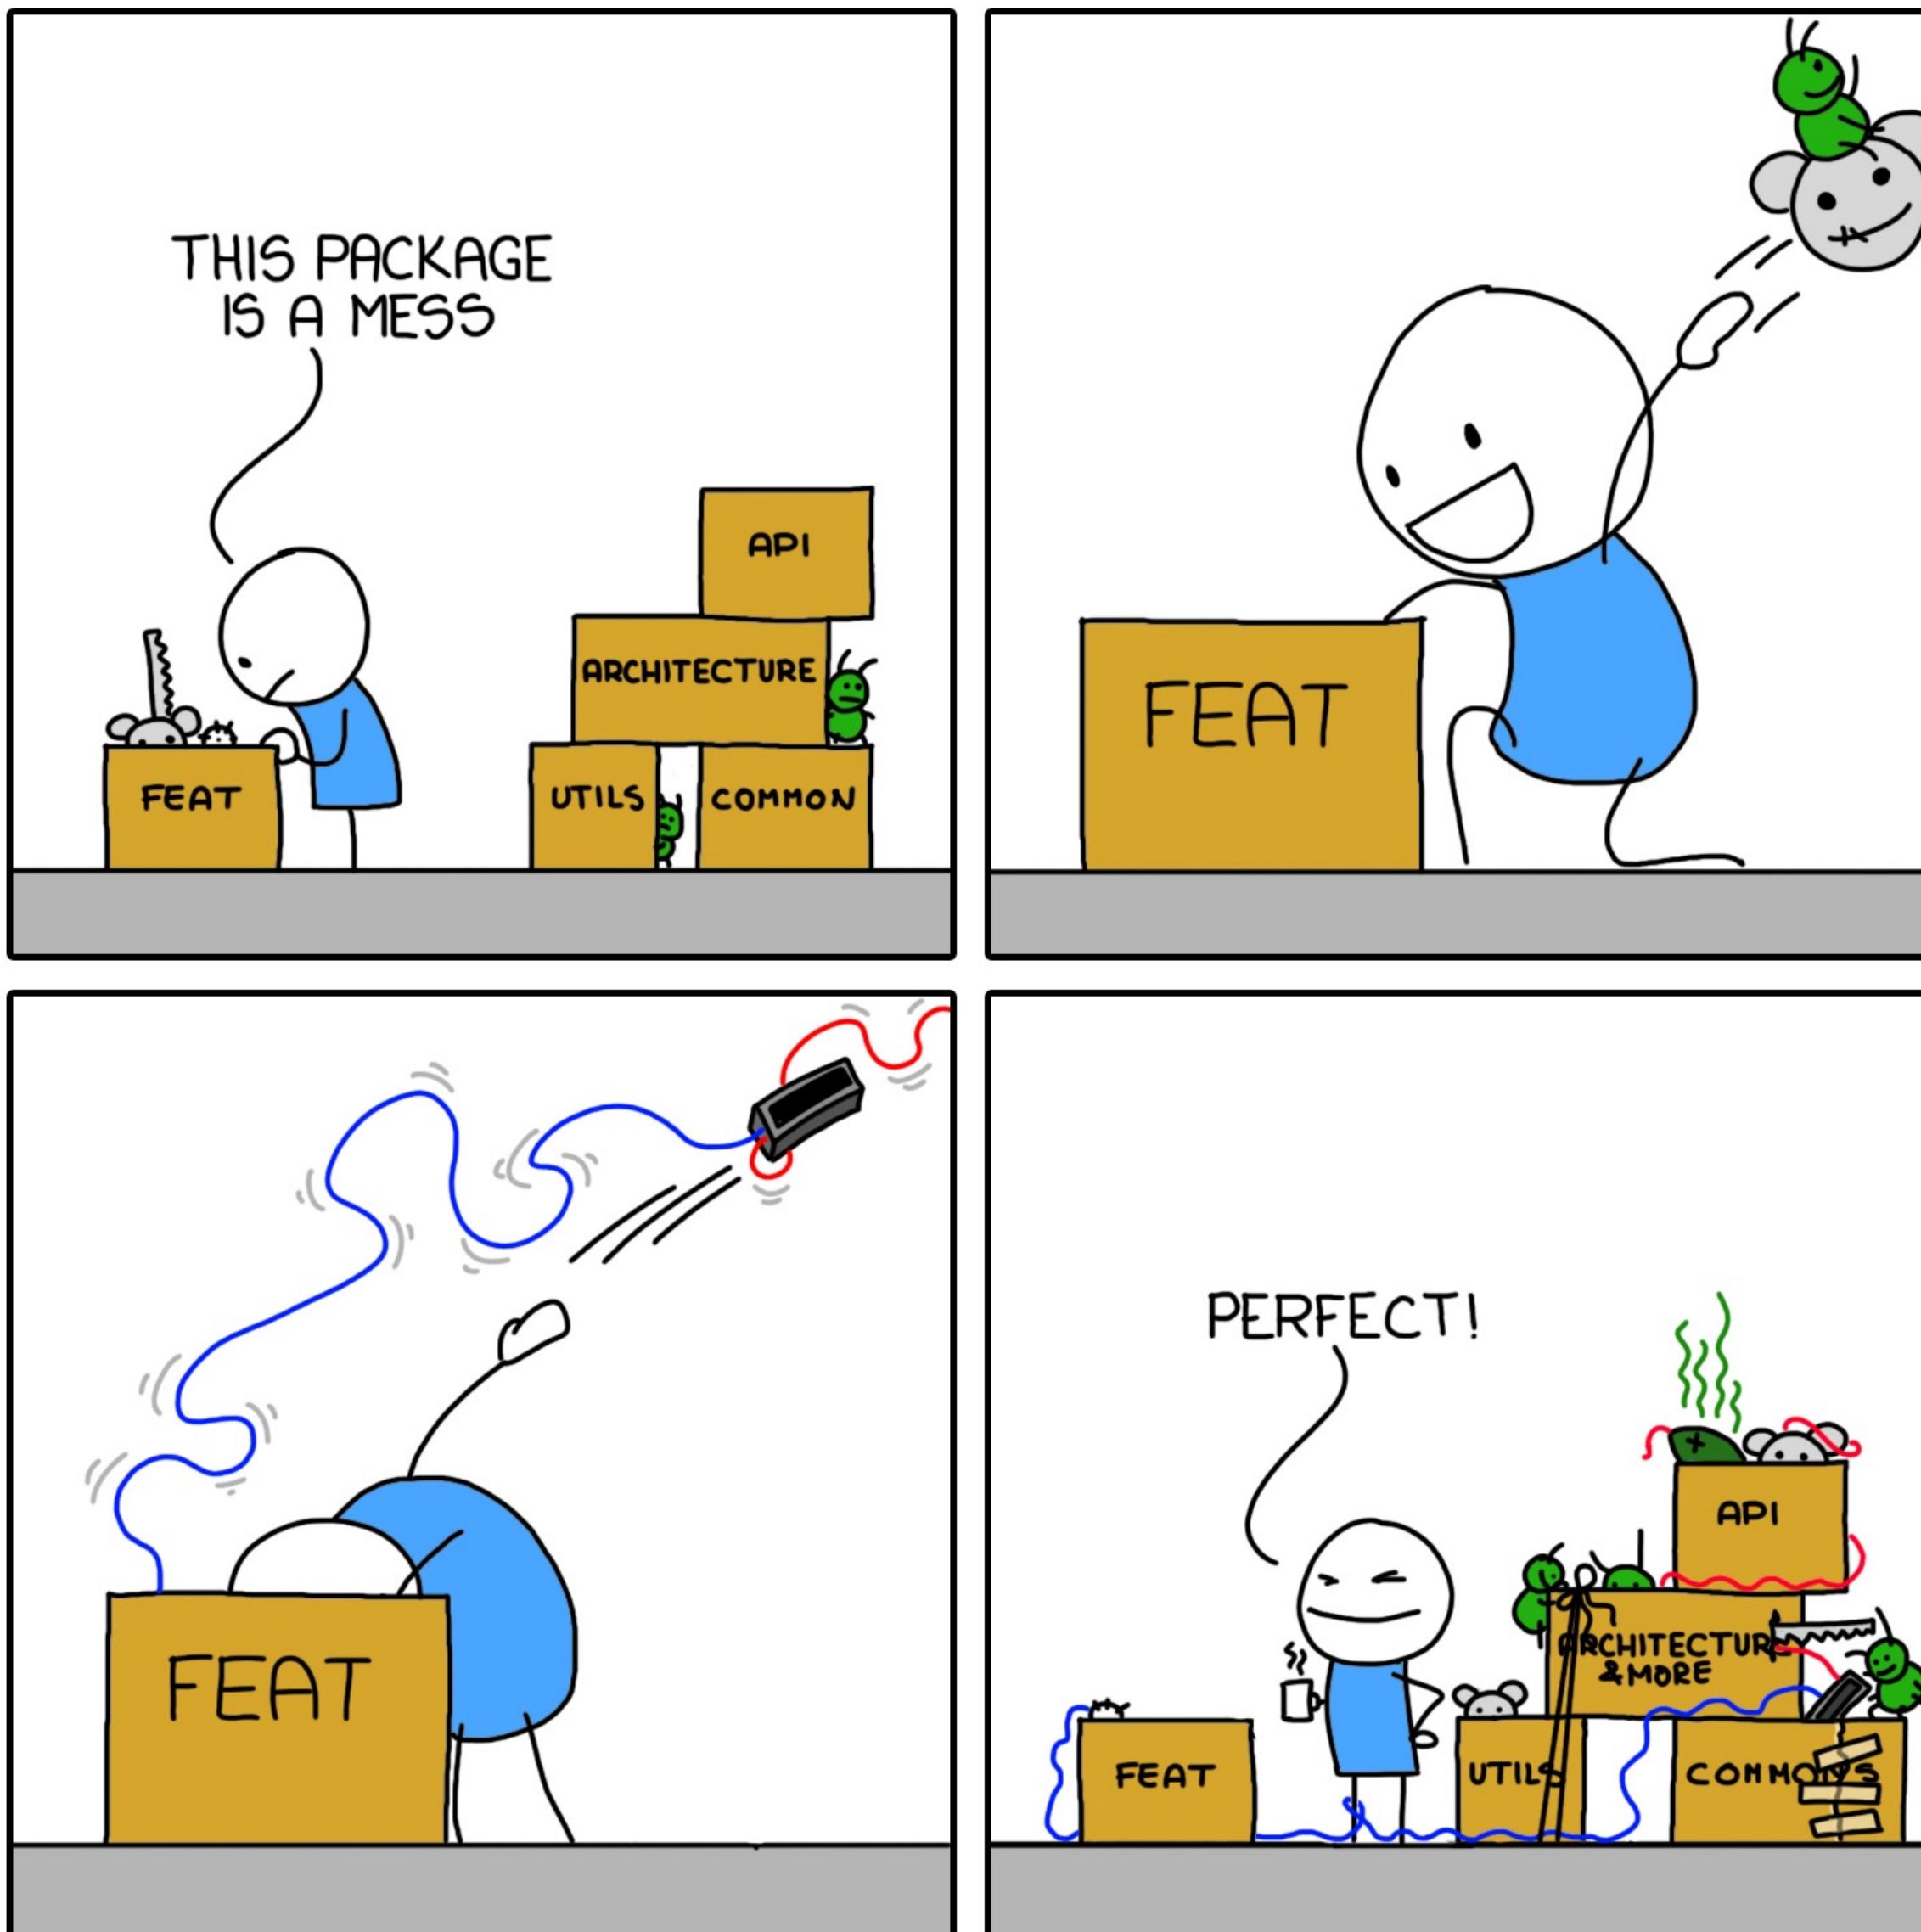
```
1010101101010101010101011001101010010011110110110110  
110101011011010101101010101010101010101100110100100  
1111011011011010101101011010101101001010101010110  
01101010010011110110110101010110101101011010110101010  
10101010110011010101001001111011011010110101101101010  
101101001010101010101100110101010010011110110110110  
1010110101010110101010101010101100110100100100111  
1011011011010101101011010101010101010101010101100110  
1010010011110110110101011010101101010110101011010101  
0101010101100110101001001111011011010110101011011011  
010101101010101010101011001101010101010101010011110110  
010110101010101010101010101010101010101010101010110110  
10101101101010110101010101010101010101010101010100111  
10110110110101101011010101010101010101010101010101010  
1100110100100111101101101011011010110110101101011010  
01010101011001101010010011110110110101101011010101010  
0101010101100110101001001111011011010101101011011010  
010110101010101010101100110100100111101101101011011010  
1010110110101101010110101010101010101010101010101010  
0100111101101101010110101101010110101010101010101010  
1100110101001001111011011010110101101011010101101010  
01101010101010101100110101010101001001111011011011010
```

# Accelerate Serverless Java

## Improved Code

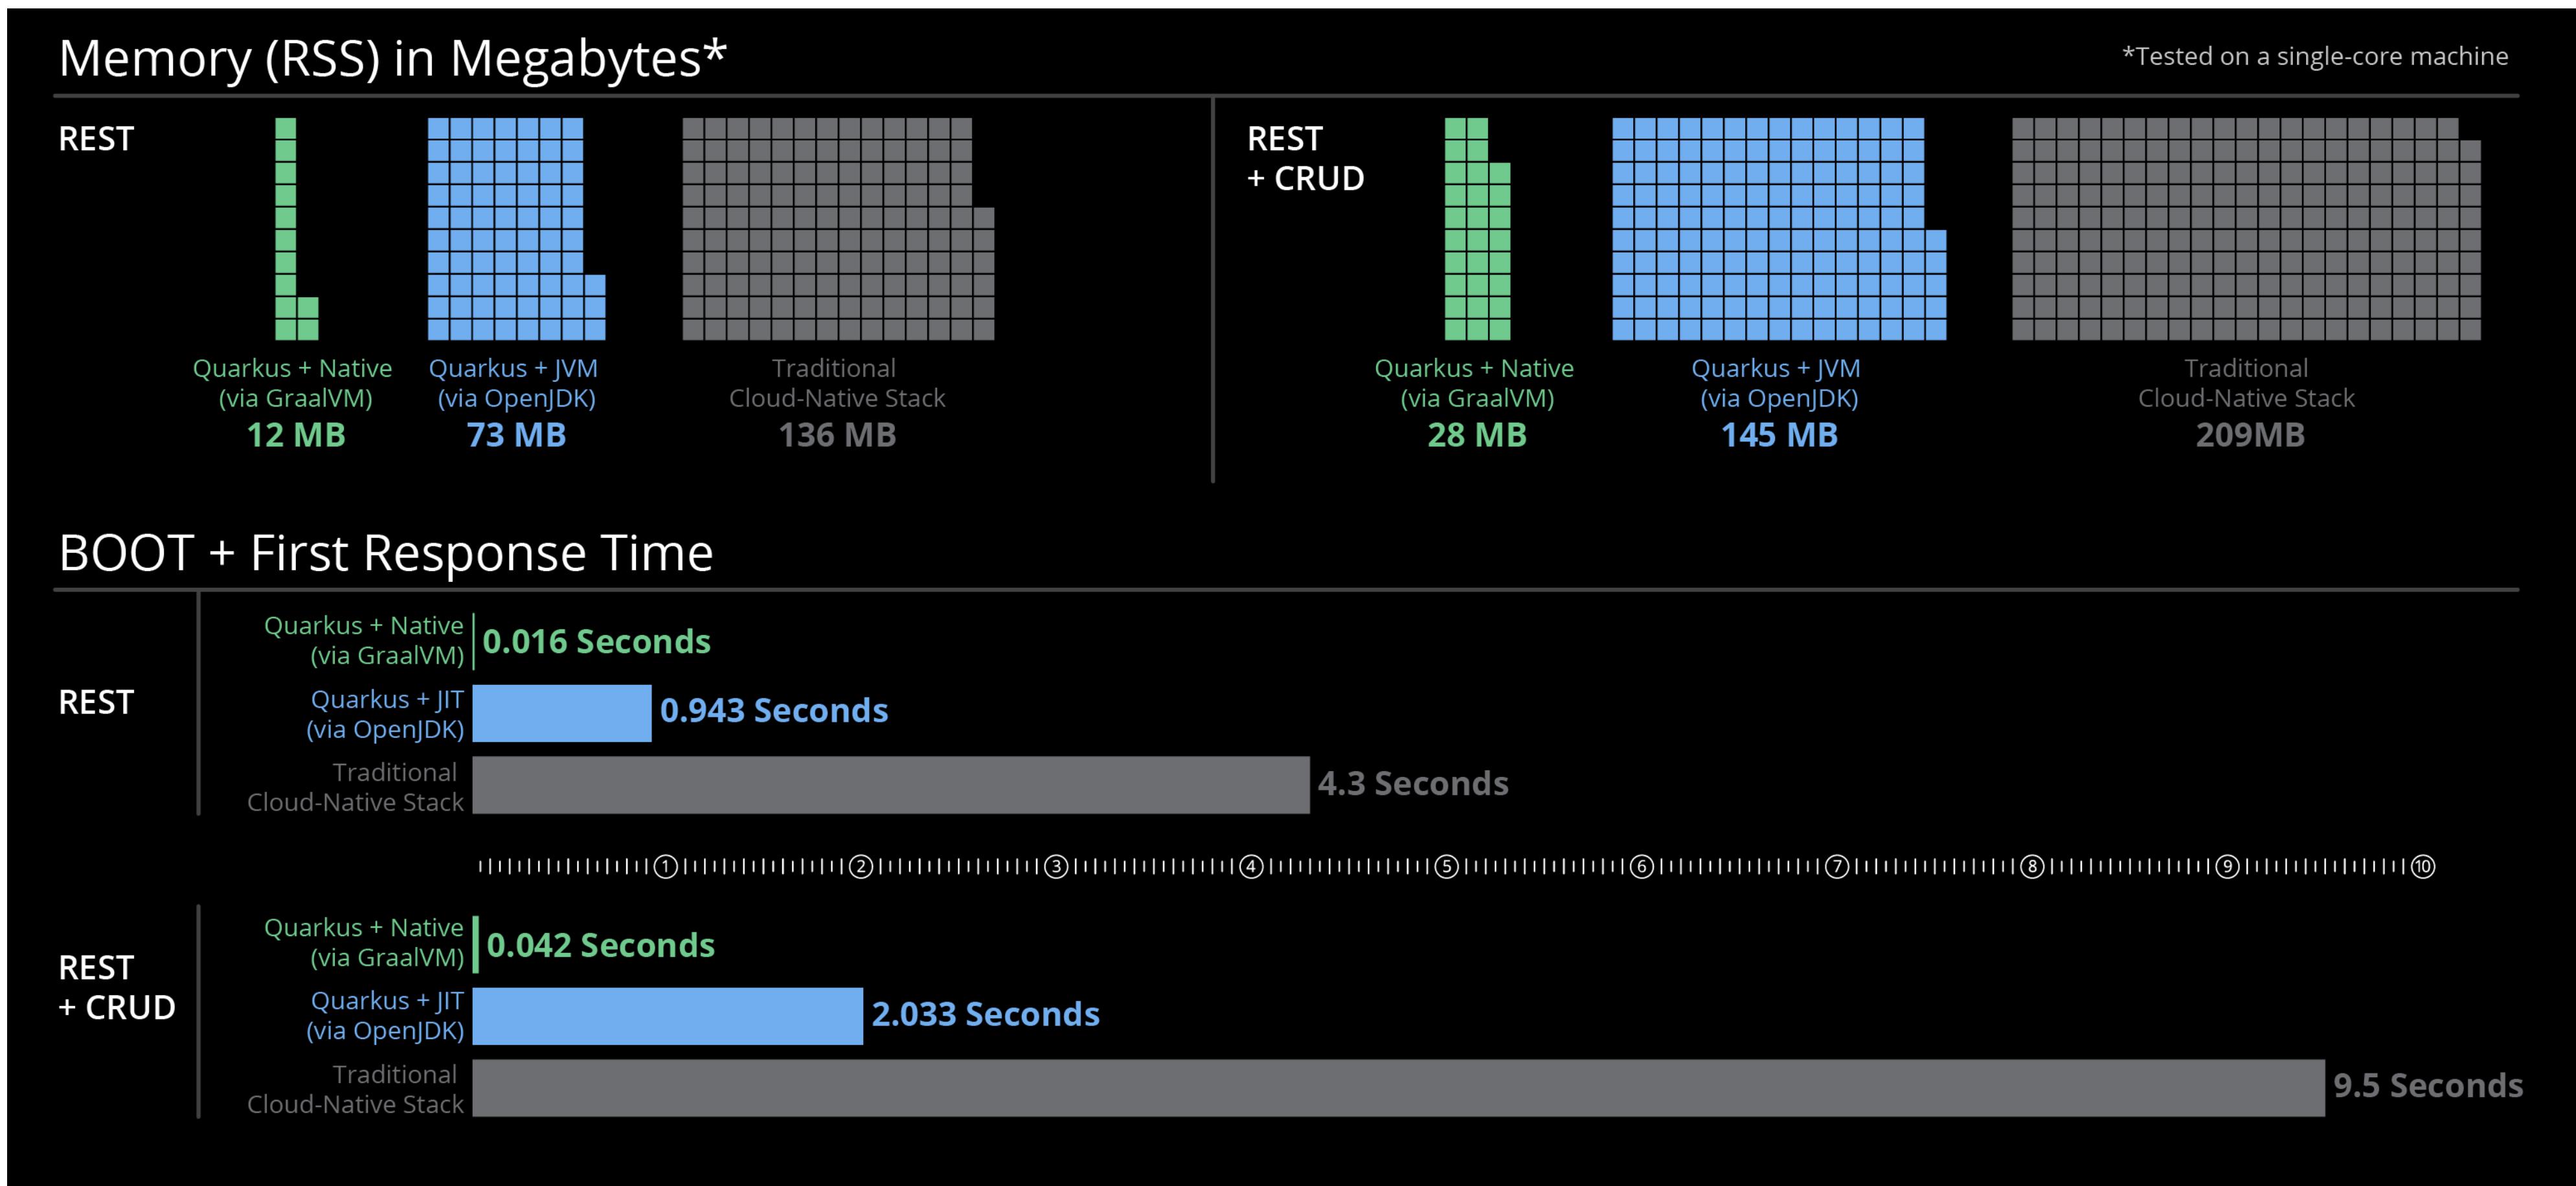
Write cloud-native Java application with serverless mindset.

1. Utilize CPU burst on startup (move everything to static) – decreased billed execution time but sacrificed performance during cold start
2. Use Compile Time frameworks such as Micronaut, Quarkus
3. Avoid reflections
4. Reduce dependencies as much as possible; Use lightweight dependencies;
5. ~25-35% improvement



# Accelerate Serverless Java

## AWS Lambda Custom Runtime + GraalVM

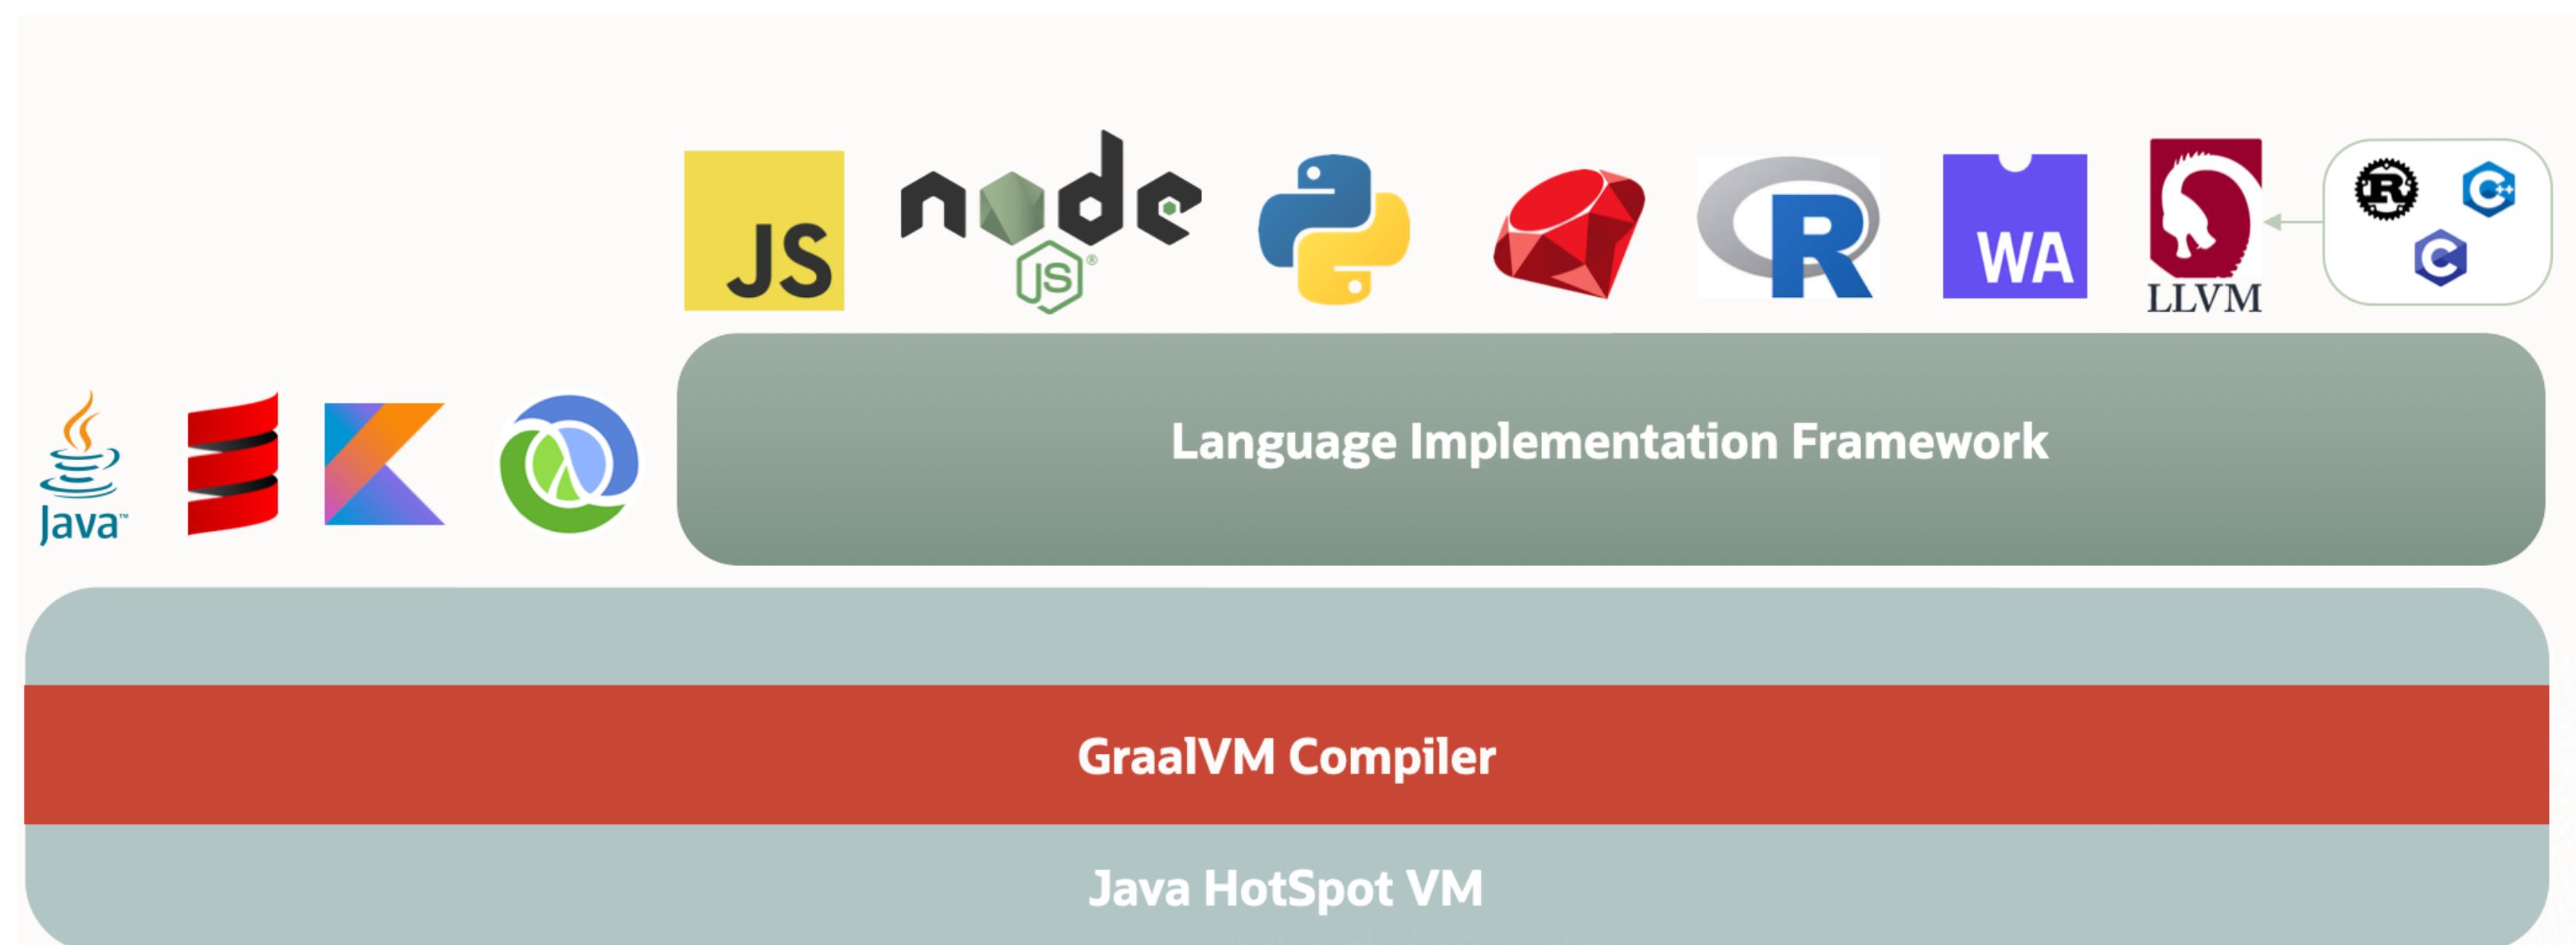


# Accelerate Serverless Java

## AWS Lambda Custom Runtime + GraalVM GraalVM

GraalVM is an umbrella term that basically consists of three different technologies

1. Graal is a high-performance **JIT compiler** written in Java.
2. **Truffle**: GraalVM's language implementation framework – which makes GraalVM a polyglot virtual machine (supporting a large set of languages: JavaScript, Ruby, Python, R, C/C++)
3. **Native Image** is a technology to compile Java code ahead-of-time to a binary – a native executable.

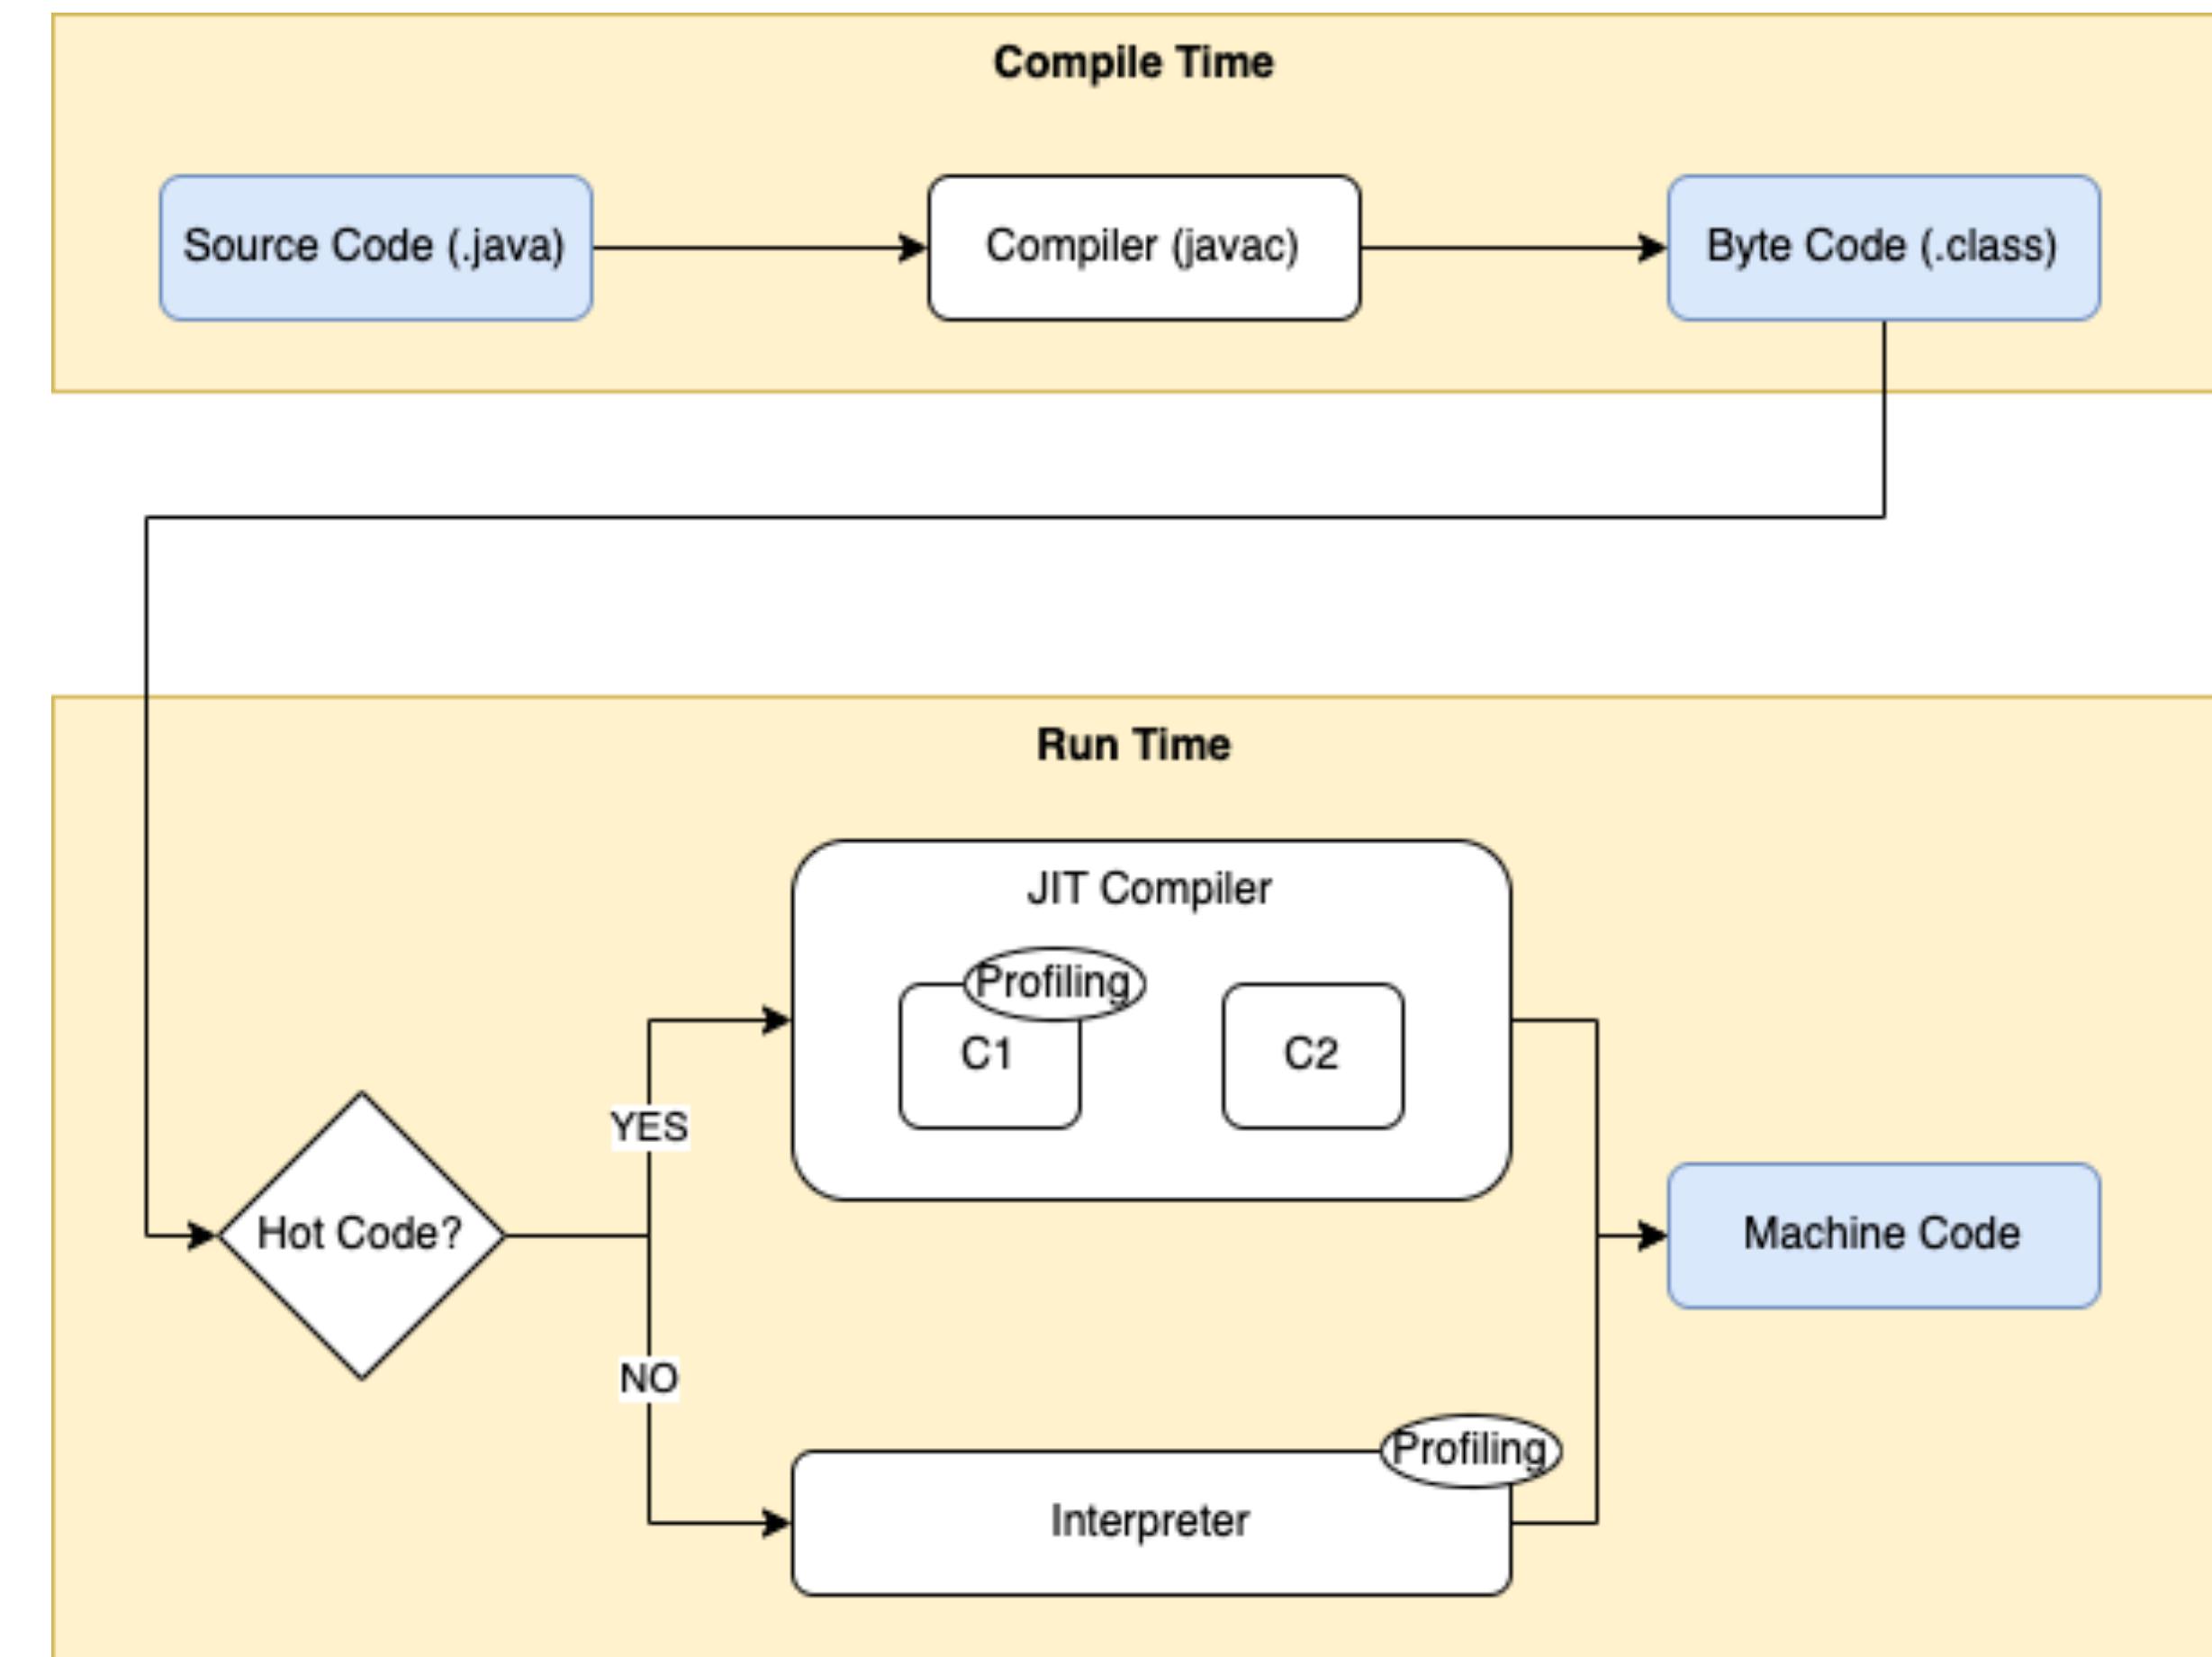
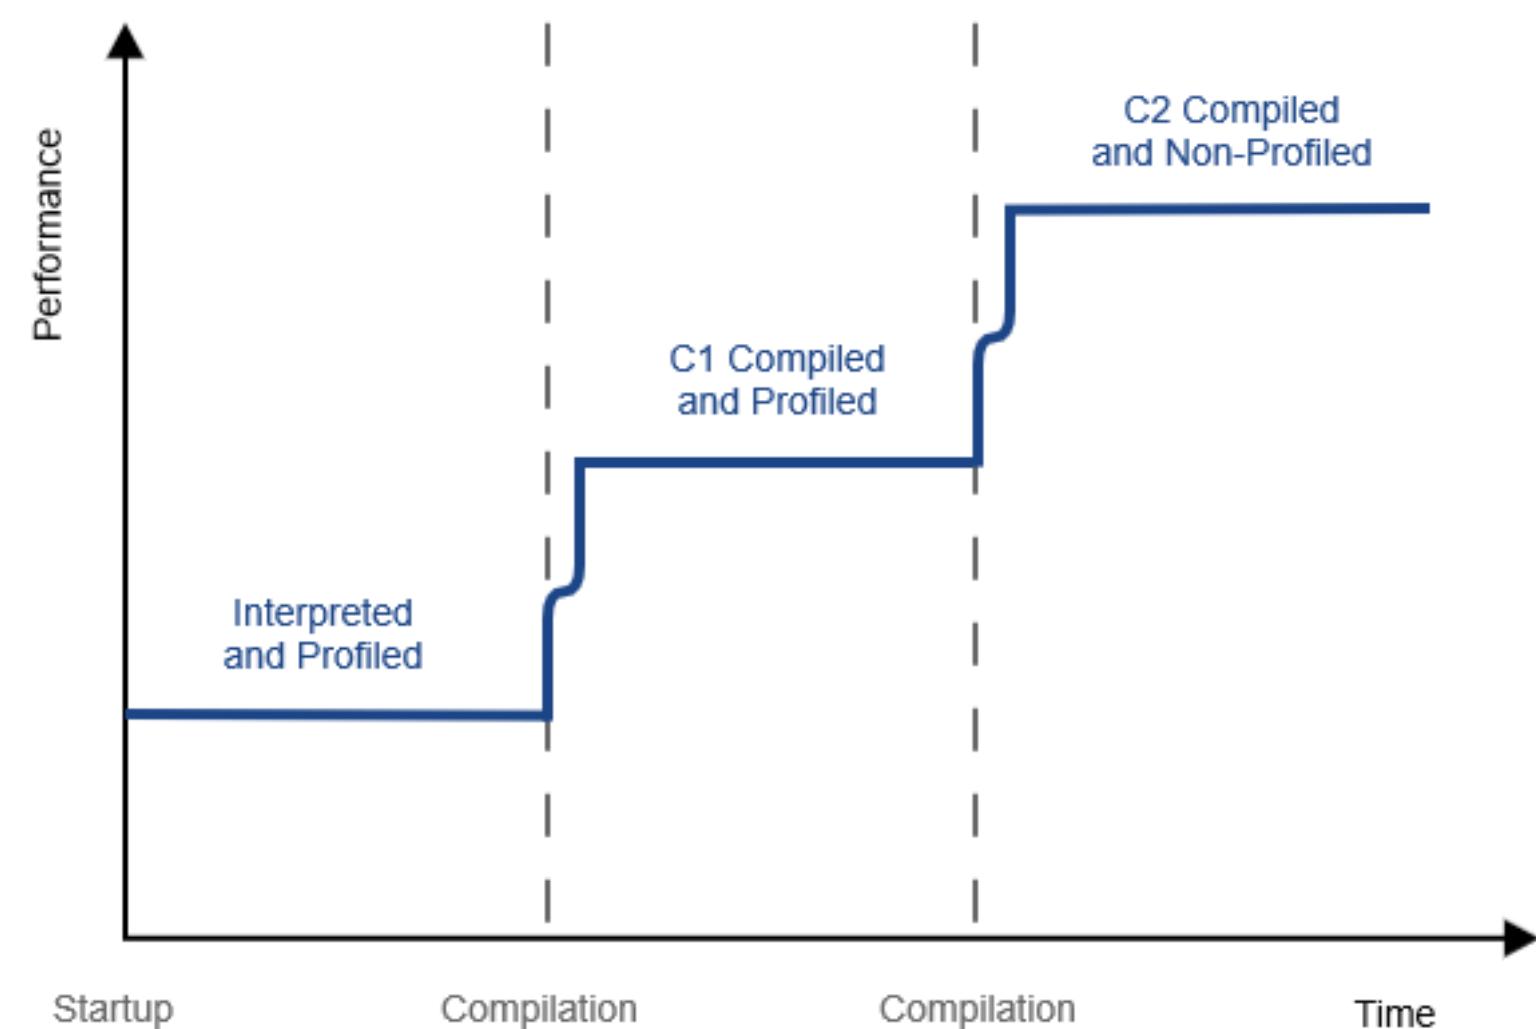


# Accelerate Serverless Java

## AWS Lambda Custom Runtime + GraalVM JIT Compilation

The JVM [interprets](#) and executes [bytecode](#) at runtime. In addition, it makes use of the just-in-time (JIT) compilation to boost performance.

A JIT compiler **compiles bytecode to native code for frequently executed sections**.



# Accelerate Serverless Java

## AWS Lambda Custom Runtime + GraalVM Native Compilation

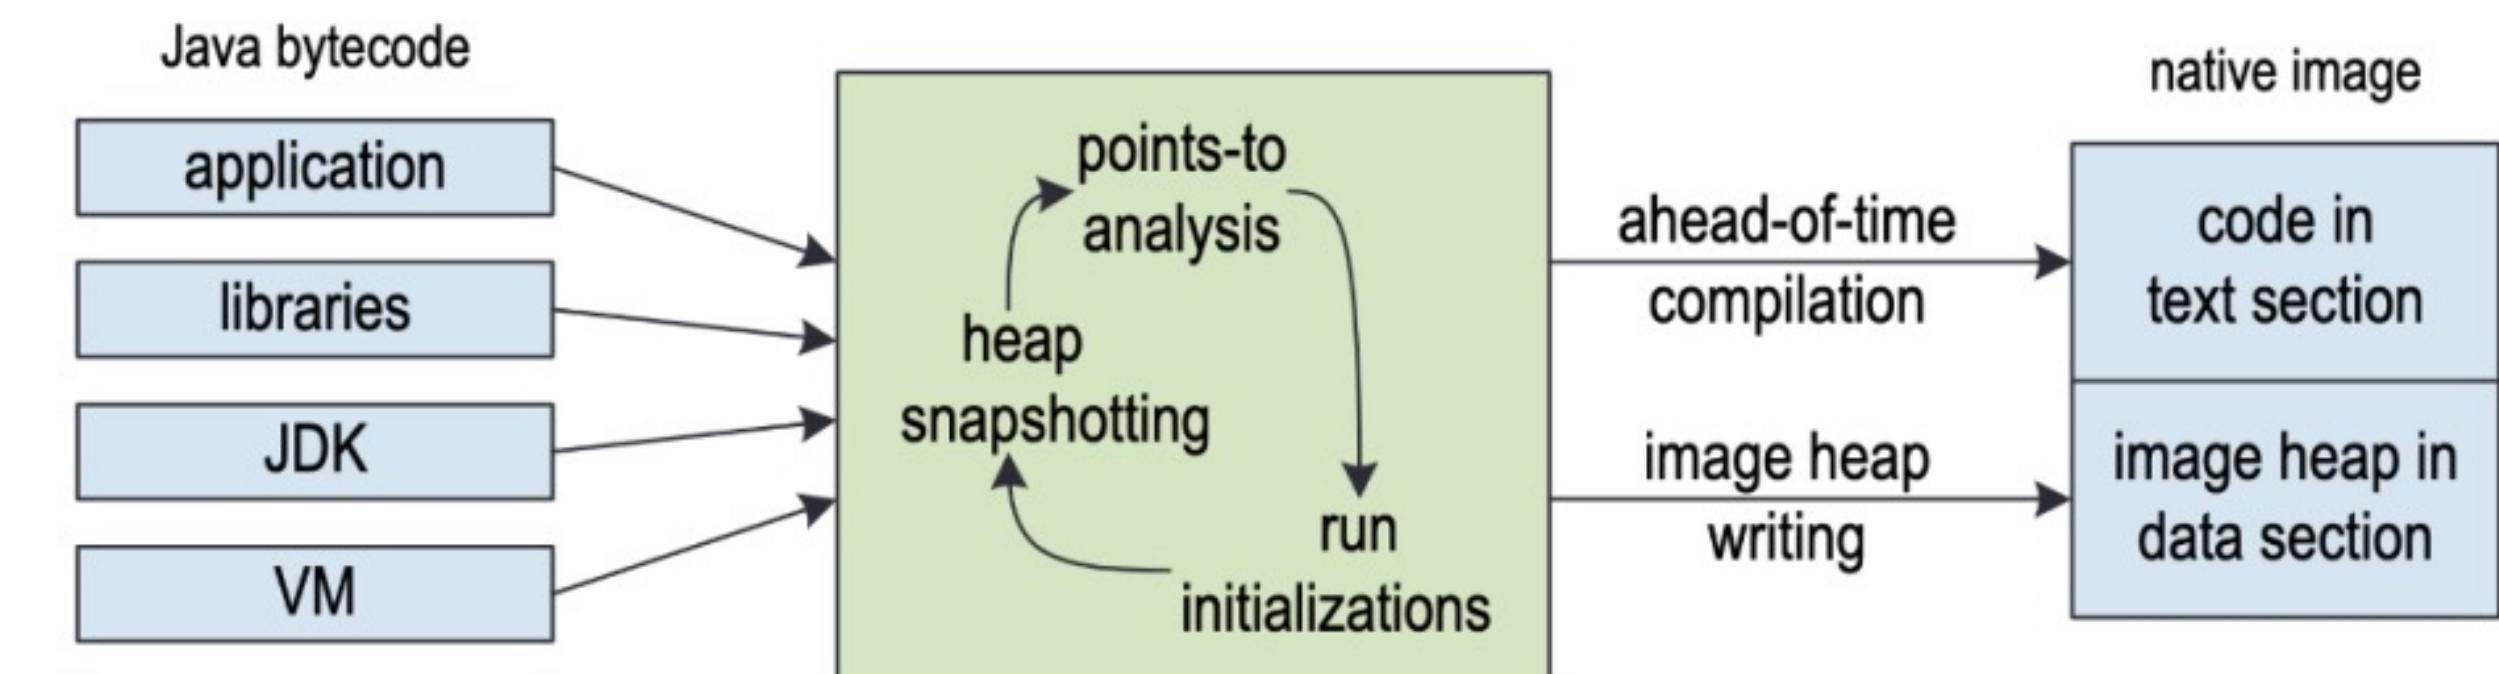
- Graal can run in both just-in-time and **ahead-of-time** compilation modes to produce native code.
- It allows statically compiling bytecode directly into machine code at **build time**

### Pros

- Decreased artifact size: use only a fraction of the resources required by the JVM.
- Applications start in milliseconds - deliver peak performance immediately, no warmup.
- Reduced attack surface (No dynamic loading, close world)

### Cons

- Close World Assumption (*solution: Tracking Agent*)
- Significantly (!) increased build time
- Platform dependent artifact



# Accelerate Serverless Java

## SnapStart

1. Announced at AWS re:Invent 2022 (**Dec 2022**)
2. Introduced **explicitly for Java runtime** in order to solve Java Cold Start issue out of the box
3. SnapStart helps you **improve startup performance** by up to **10x at no extra cost**
4. **Snapshot** created from the initialized function at deploy time. (i.e.: state after the cold start)
5. Lambda resumes the function from the cached snapshot upon invocation

### Publish Version (SnapStart)



```
--snap-start ApplyOn=PublishedVersions
```

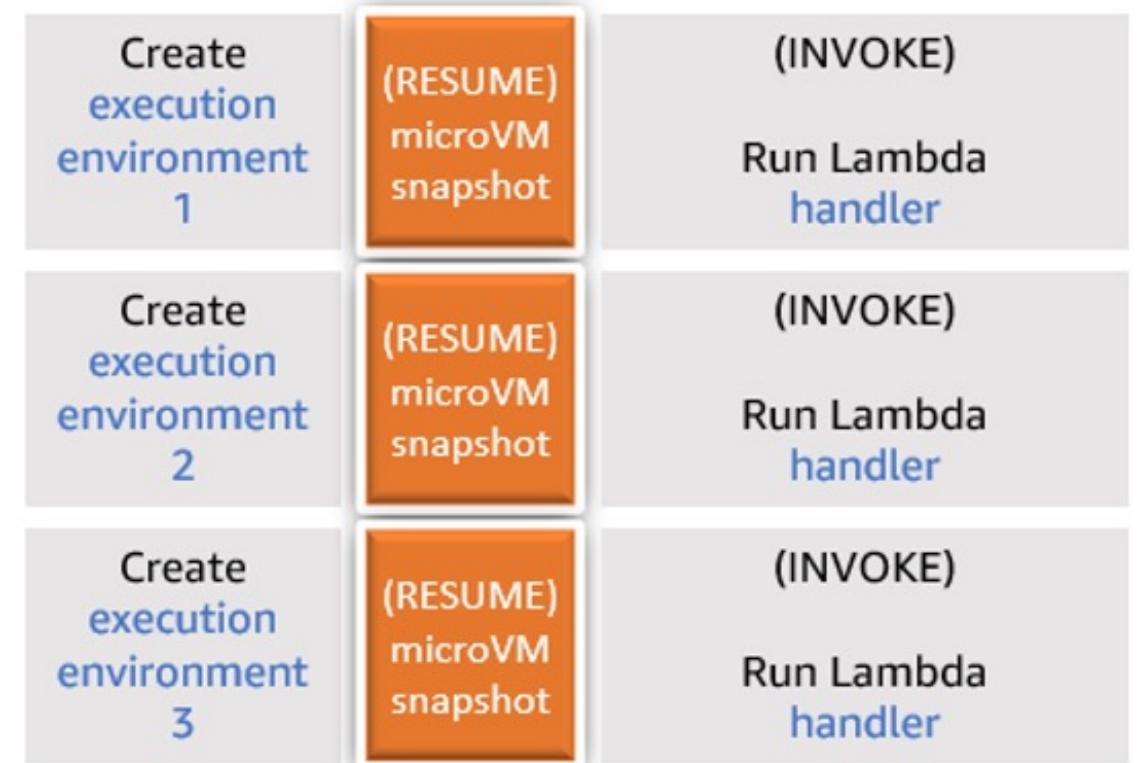
```
publish-version
```

<https://docs.aws.amazon.com/lambda/latest/dg/snapstart.html>

<https://catalog.workshops.aws/java-on-aws-lambda/en-US/01-migratio>

<https://quarkus.io/blog/quarkus-support-for-aws-lambda-snapstart/>

### Request Lifecycle (SnapStart)



# Accelerate Serverless Java

## Comparison

Optimization	p0	p95	Improvement in %	SnapStart p0
No optimization	11.378	13.256		1.287
Tiered compilation	6.725	8.387	~40%	1.15
Lightweight dependencies	5.052	6.245	~25%	0.901
Function handler (Spring Cloud functions)	4.571	5.606	~10%	0.779
Micronaut JVM	3.499	4.284	~25%	0.765
No framework	1.850	2.588	~45%	0.705
GraalVM (Spring Native)	0.683	0.890	~65%	n/a

<https://catalog.workshops.aws/java-on-aws-lambda/en-US/01-migration/results>

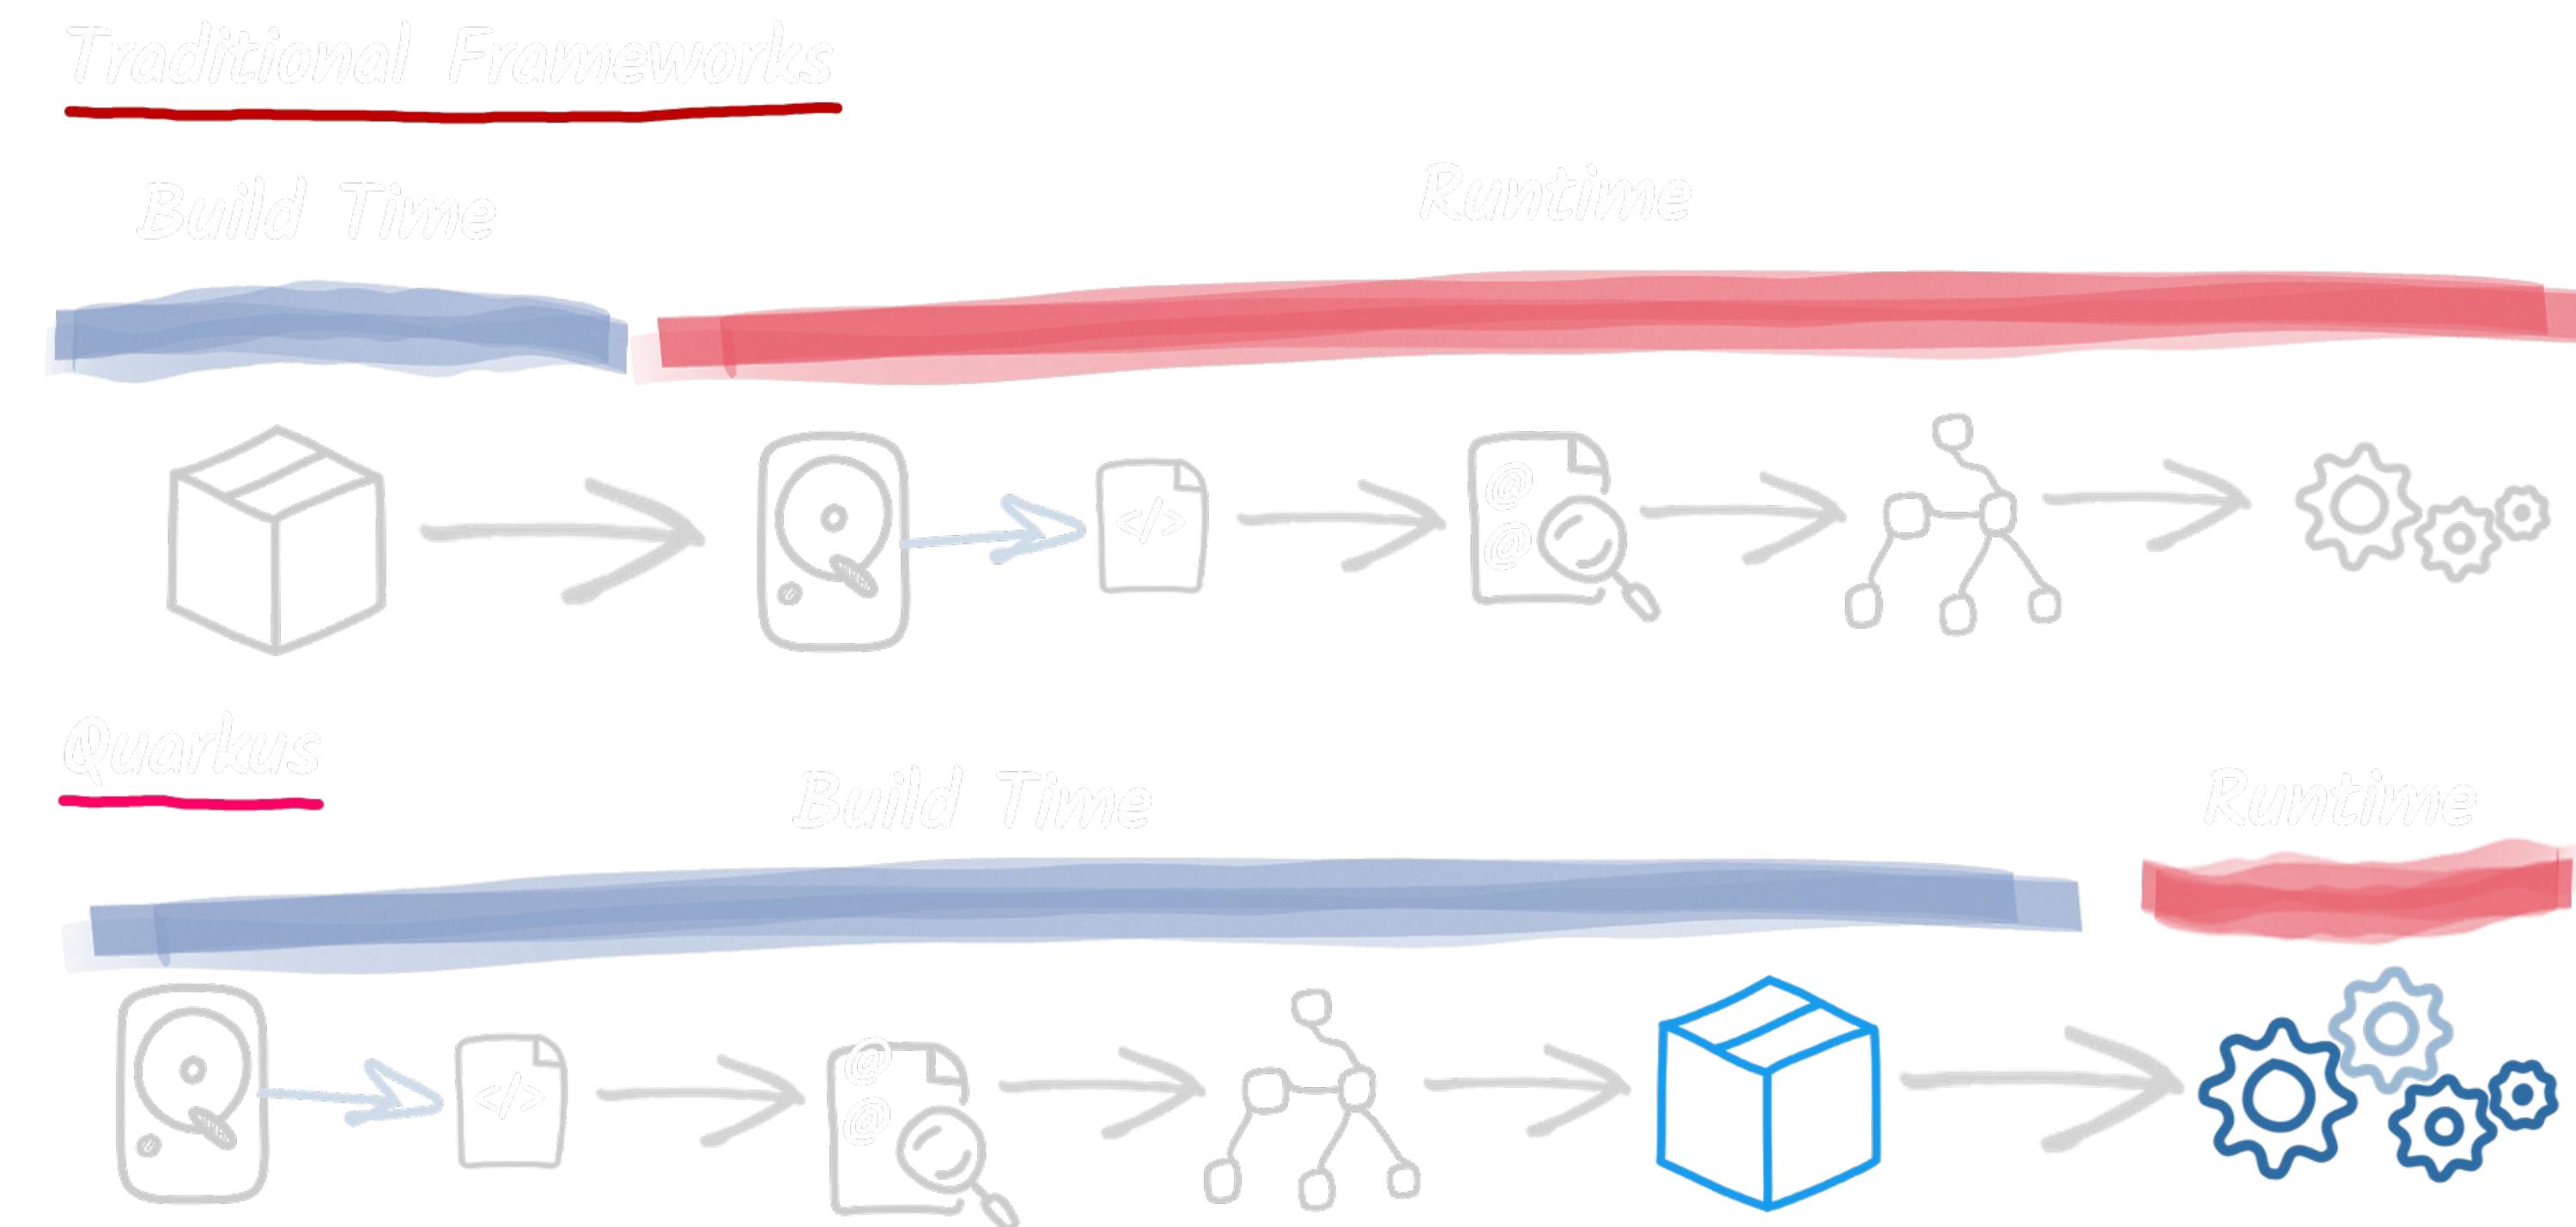
# Demo #3. Enabling SnapSt





# Compile Time Framework – Shift left

Both frameworks are optimised for low memory usage and fast startup times. Do at build-time what traditional frameworks do at runtime. Generate meta-information (in form of additional byte codes) at build time about your code.



<https://quarkus.io/container-first/>

<https://micronaut.io/2023/03/21/micronaut-framework-code-generation/>

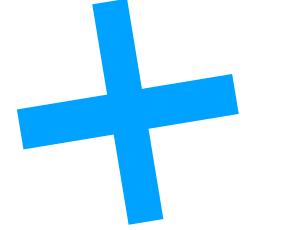
# Build Time Actions

What artifacts does **Micronaut** generate as bytecode?

1. **Bean definitions** to power the Micronaut Framework dependency injection engine.
2. **Bean Introspections** to power features such as **reflection-free** serialization.
3. **Proxies** to power AOP features.
4. **Bean Factories**. For example, Micronaut Kubernetes and Micronaut Oracle Cloud generate factories for external SDK Clients from an annotation processor.

What artifacts does **Quarkus** generate as bytecode?

1. Custom **proxies** instead of dynamic proxies
2. Replace **reflection** calls with regular invocations
3. Arc (Quarkus DI framework): eliminates reflections, **evaluates DI injection graph at build time**



GraalVM Native Executable support has been an essential part of the design for Quarkus from the beginning

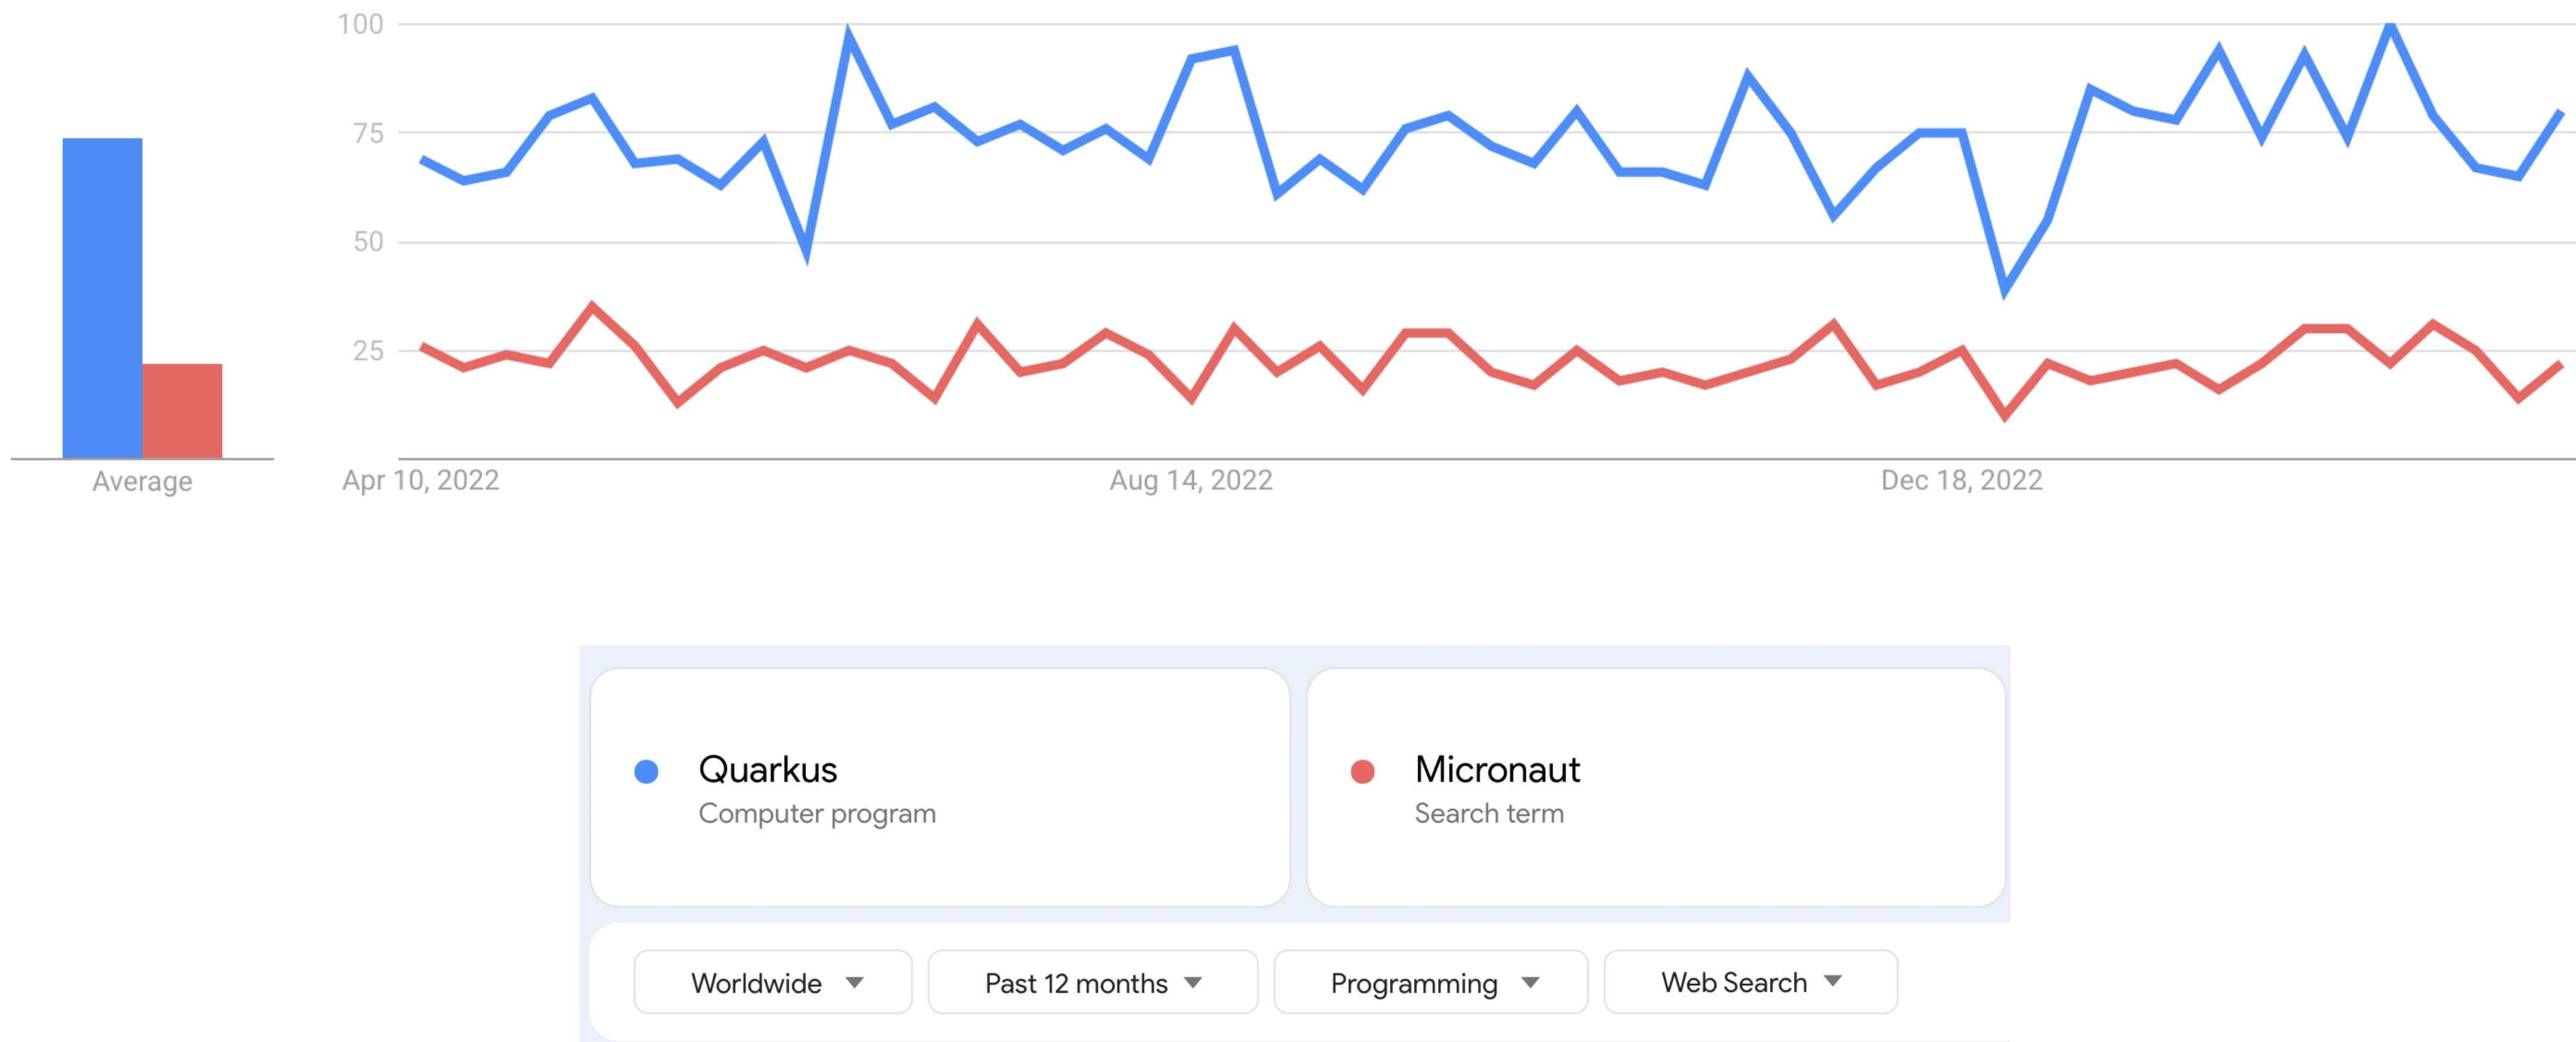
<https://quarkus.io/container-first/>

<https://micronaut.io/2023/03/21/micronaut-framework-code-generation/>

<https://quarkus.io/guides/building-my-first-extension>

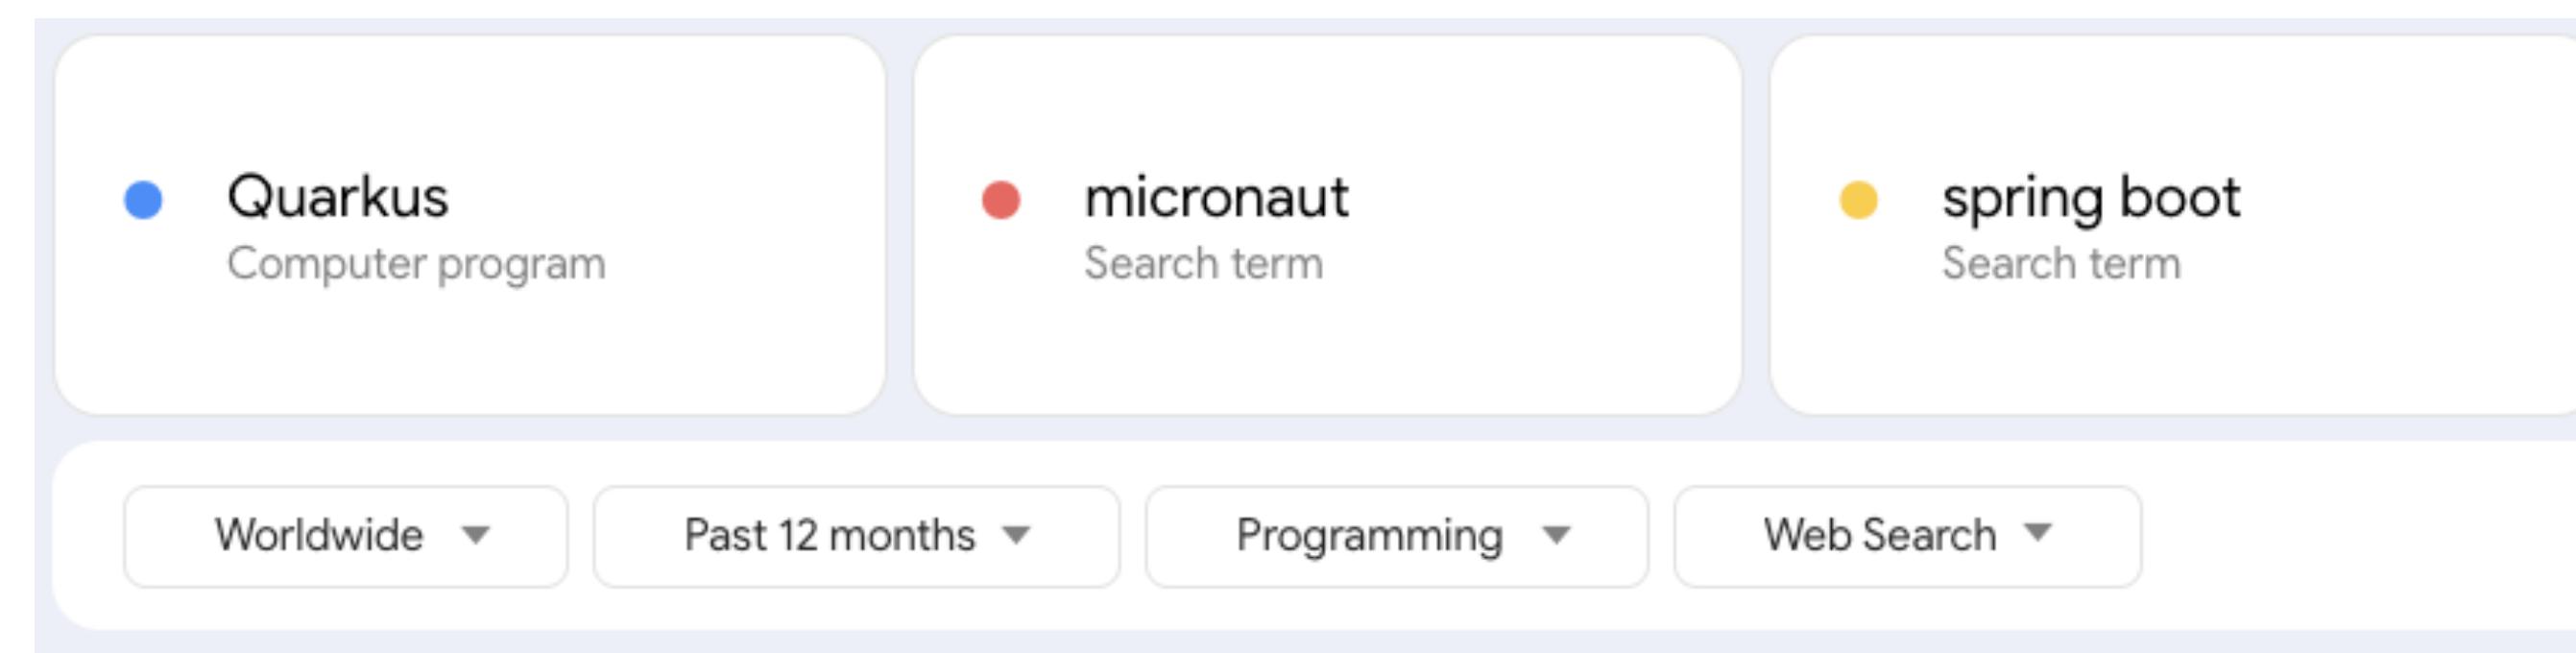
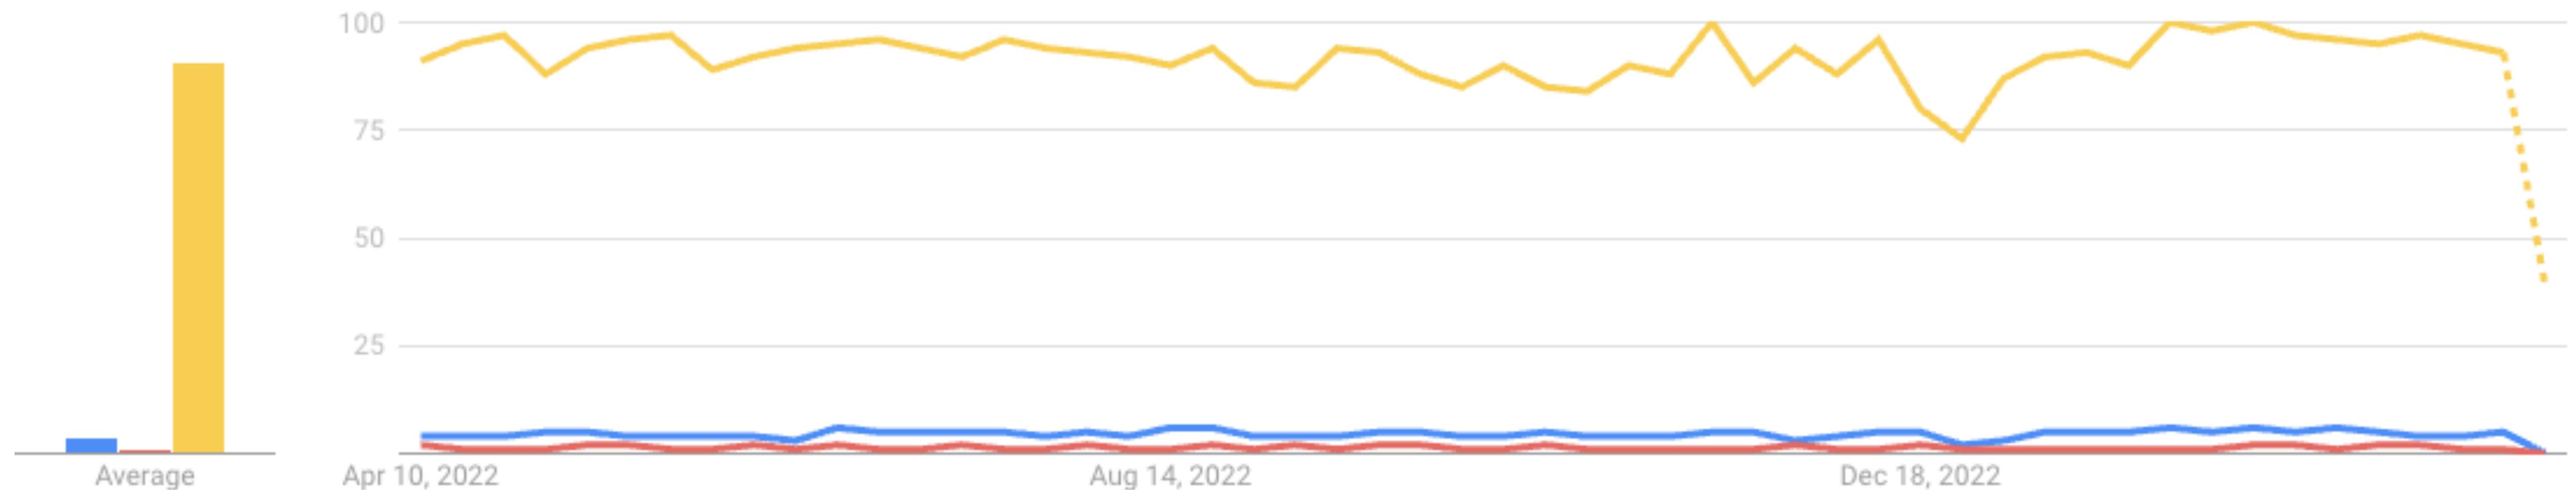
# Comparison

## Google Trends



# Comparison

## Google Trends / Part II



# Comparison

## Github

### Micronaut

1. **Nr of Releases:** 154
2. **Contributors:** 365
3. **Stars:** 5700
4. **Forks:** 971
5. **License:** Apache 2.0

x2

*Initial Release Date: May 2018*

<https://github.com/micronaut-projects/micronaut-core>

<https://github.com/quarkusio/quarkus>

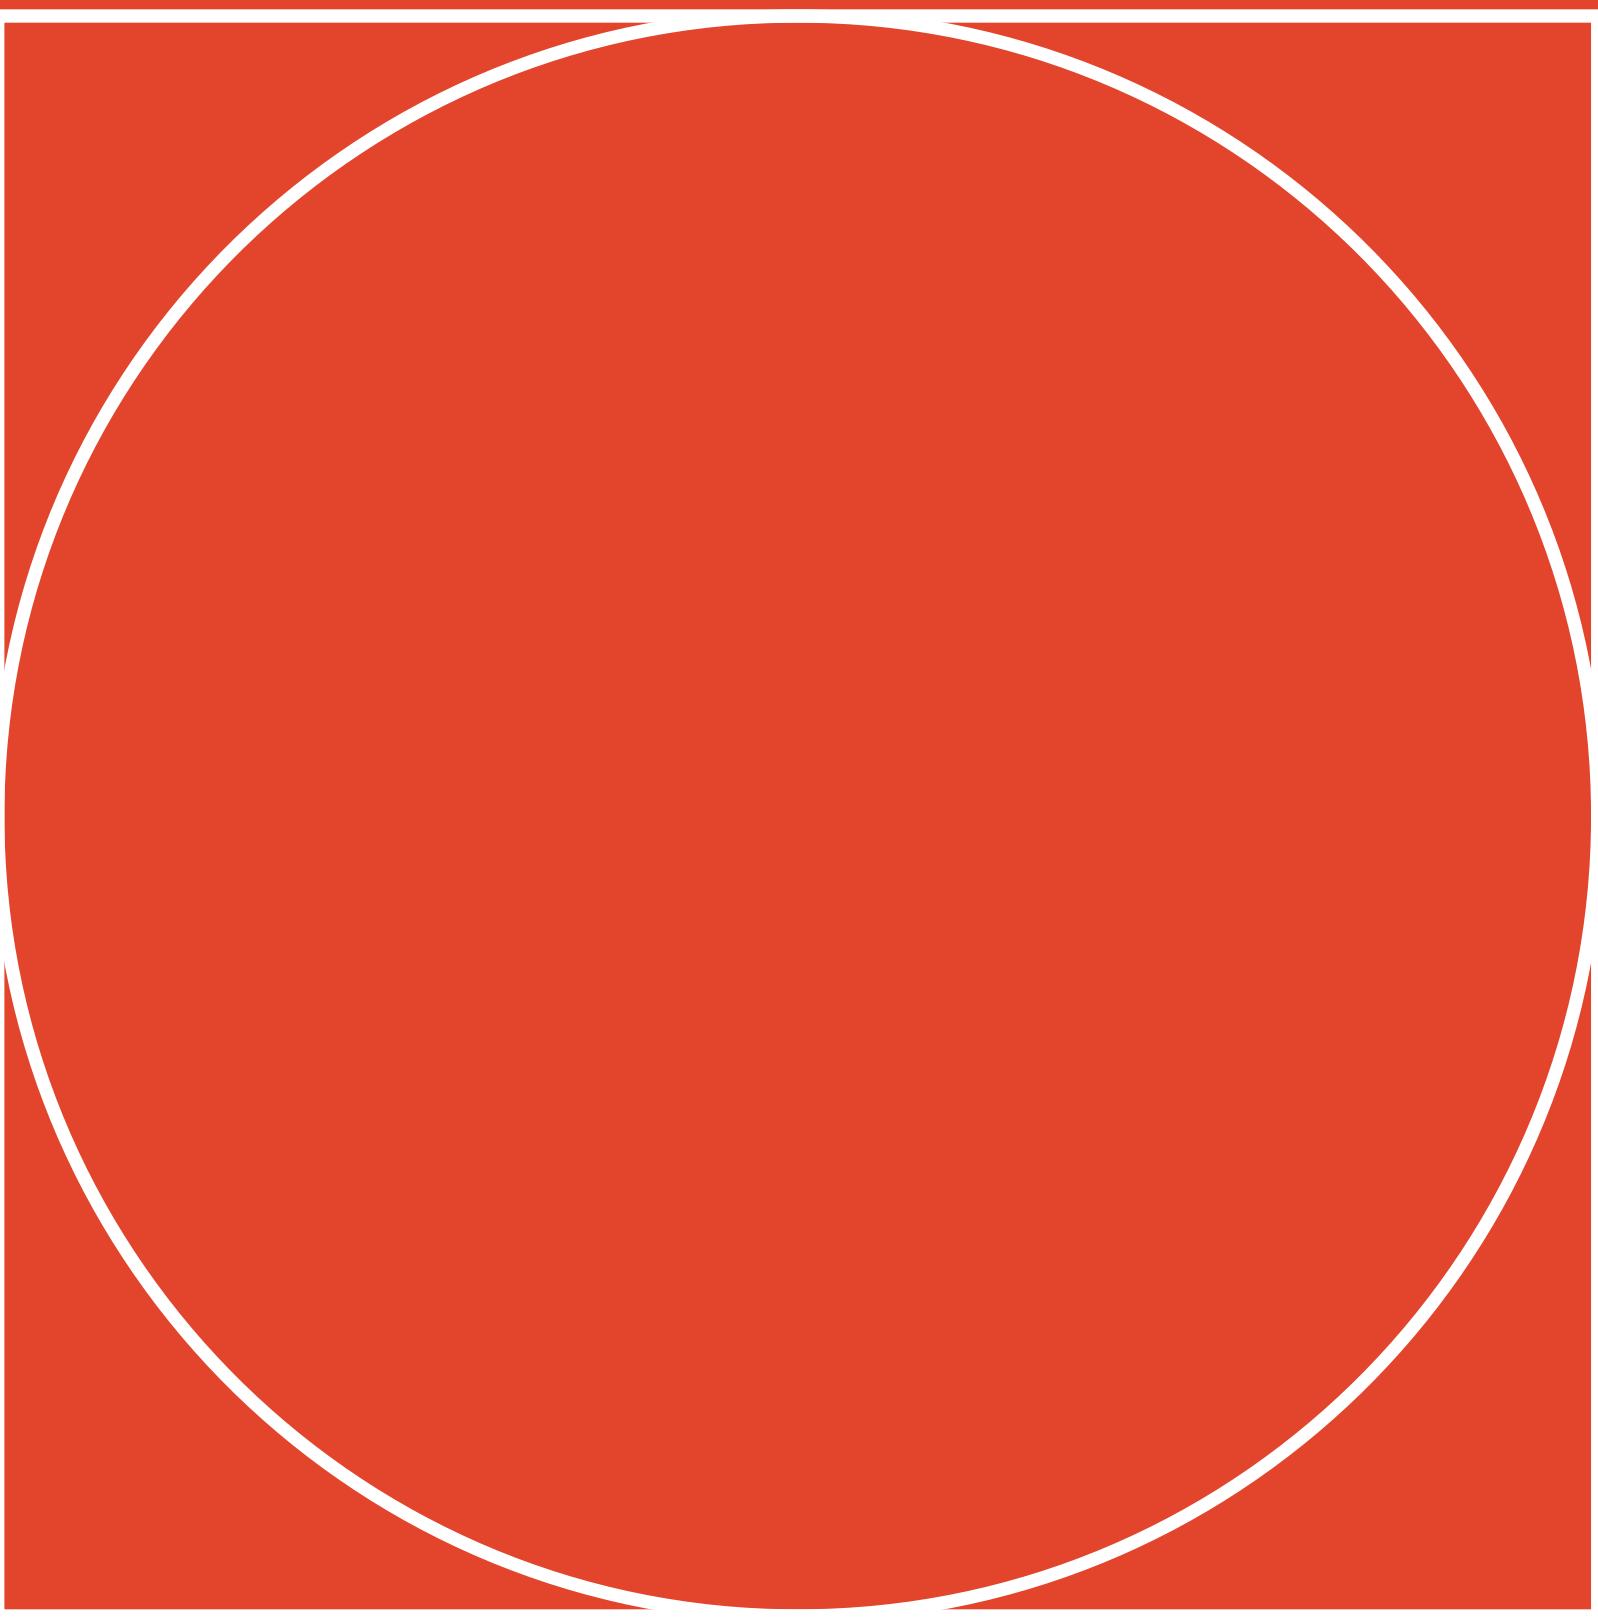
### Quarkus

1. **Nr of Releases:** 226
2. **Contributors:** 785
3. **Stars:** 11600
4. **Forks:** 2200
5. **License:** Apache 2.0

*Initial Release Date: March 2019*

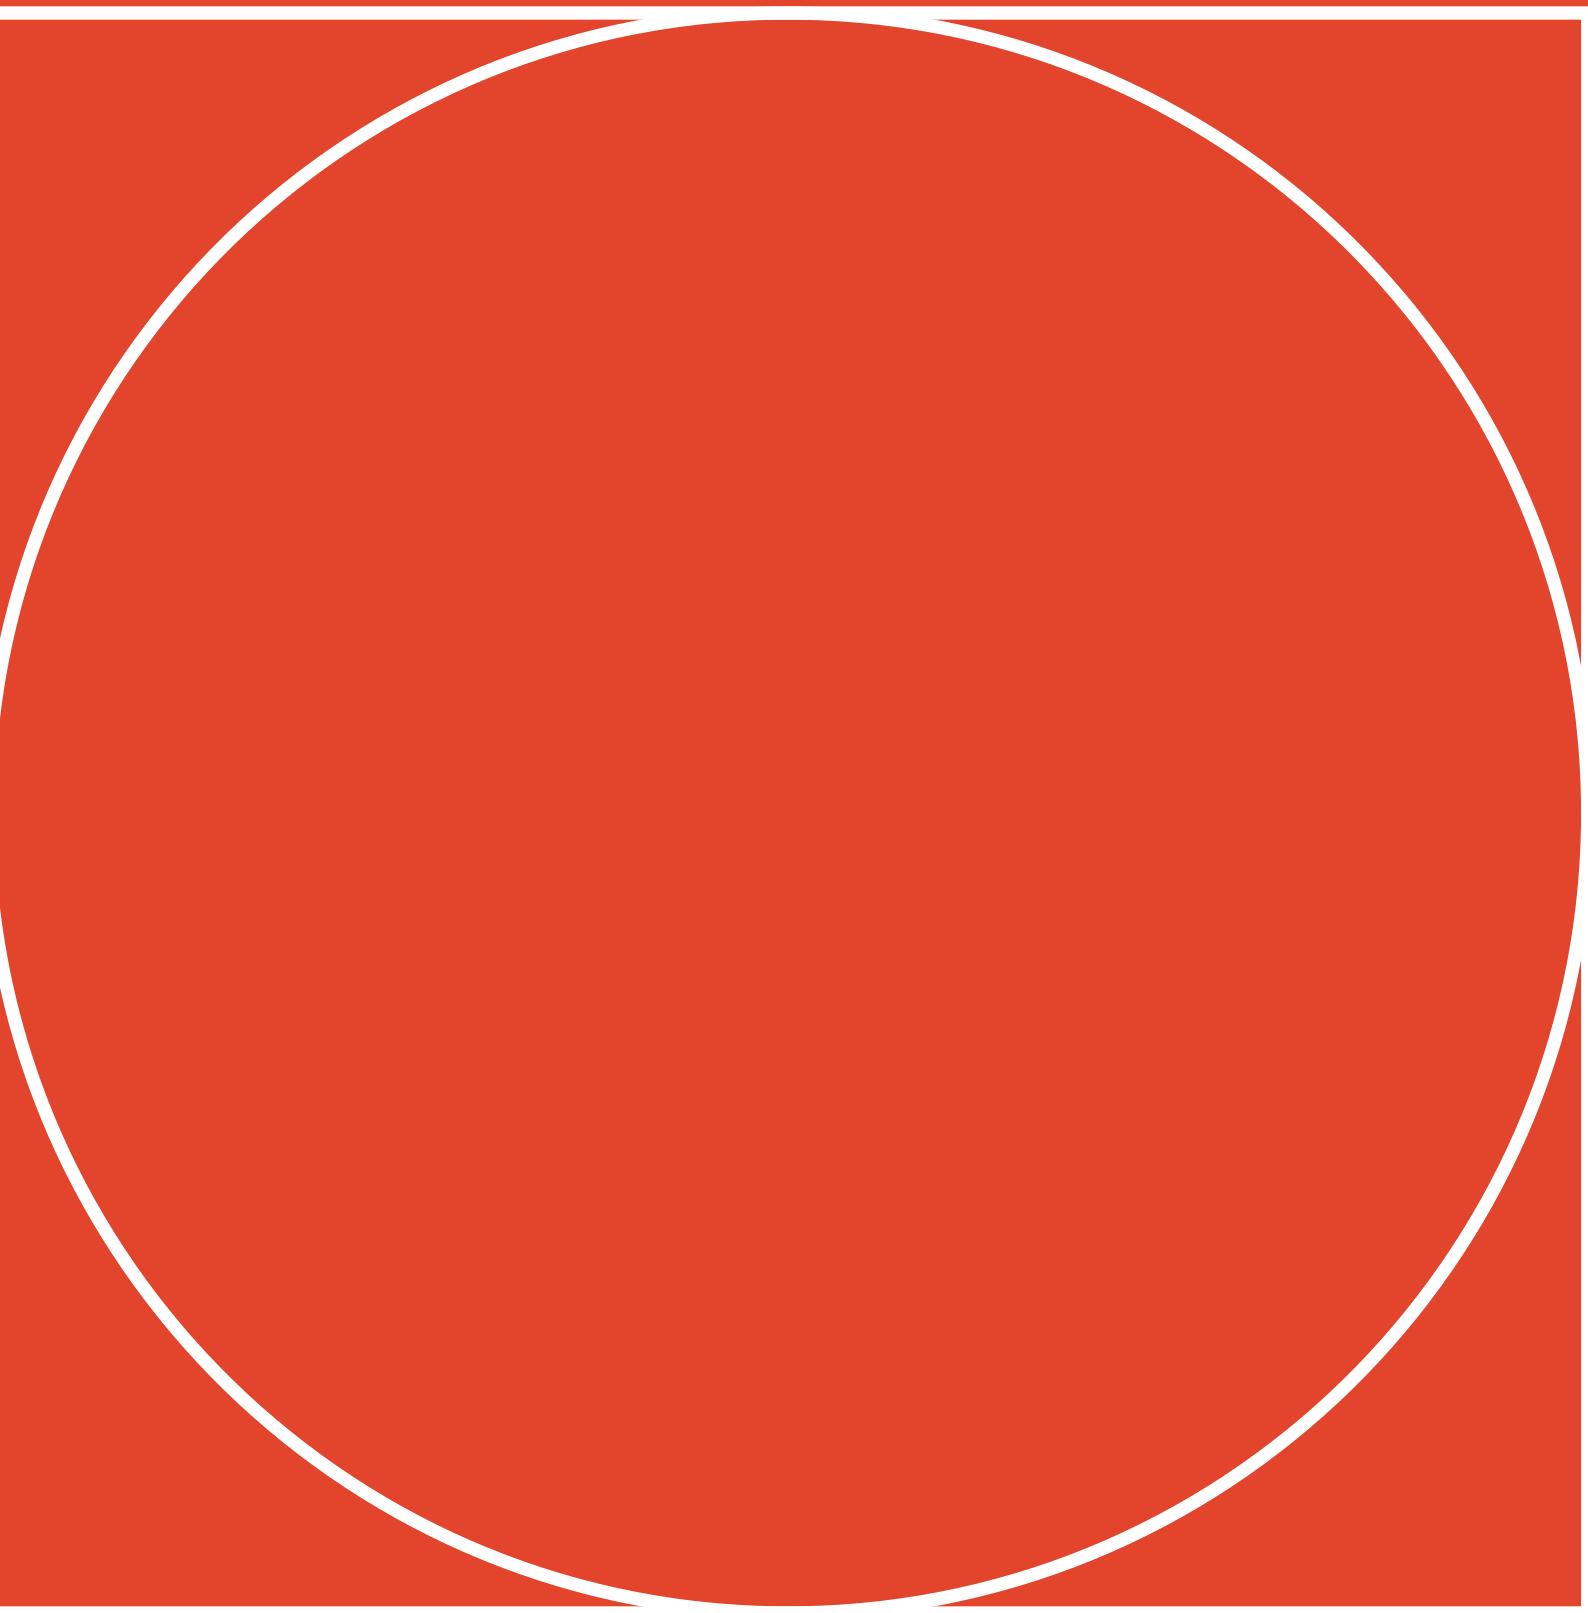
# Demo #4.1

## AWS Lambda w/ Quarkus



# Demo #4.2

## AWS Lambda w/ Micronaut



# Performance Comparison of the Hello World Lambda

## Micronaut

Duration: 395.65 ms Billed Duration: 810 ms Memory Size: 256 MB Max Memory Used: 91 MB Init Duration: 414.21 ms  
Duration: 122.96 ms Billed Duration: 123 ms Memory Size: 256 MB Max Memory Used: 92 MB  
Duration: 132.69 ms Billed Duration: 133 ms Memory Size: 256 MB Max Memory Used: 93 MB  
Duration: 392.77 ms Billed Duration: 393 ms Memory Size: 256 MB Max Memory Used: 93 MB

## Quarkus

Duration: 607.88 ms Billed Duration: 839 ms Memory Size: 256 MB Max Memory Used: 63 MB Init Duration: 230.68 ms  
Duration: 126.73 ms Billed Duration: 127 ms Memory Size: 256 MB Max Memory Used: 64 MB  
Duration: 408.96 ms Billed Duration: 409 ms Memory Size: 256 MB Max Memory Used: 65 MB  
Duration: 125.66 ms Billed Duration: 126 ms Memory Size: 256 MB Max Memory Used: 65 MB

<https://quarkus.io/container-first/>

<https://micronaut.io/2023/03/21/micronaut-framework-code-generation/>

# What we haven't touched...

## AWS Lambda Deployment Alternative

- [Deploying lambda functions as container images \(from ECR\)](#)

## Additional Frameworks and Supports

- [Spring Boot 3 and Spring Framework 6 official native support \(since September 2022\)](#)
- [Quarkus Funky \(Cloud agnostic Functions\)](#)

## Building Event Driven Application

- Choreography vs. Orchestration
- [Event Driven Architecture](#)
- [AWS Step Functions – Serverless workflow orchestration](#)
- [Lambda proxy integrations in API Gateway](#)
- And much more: AWS Event Bridge, SQS and SNS

THANKS

---

---

beyond

beyond